



NVIDIA Base Command Manager 11

NVIDIA Mission Control Manual

Revision: 25e5965d7

Date: Thu Dec 11 2025

Table Of Contents

Table Of Contents	3
0.1 About This Manual	7
0.2 About The Manuals In General	7
0.3 Getting Administrator-Level Support	8
0.4 Getting Professional Services	8
1 Introduction	9
1.1 NVIDIA Mission Control Requirements And Features	9
1.1.1 NVIDIA Mission Control Features	9
1.1.2 Checking That The NVIDIA Mission Control Edition Is Running	10
2 NMX Settings For NVLink Monitoring	11
2.1 Introduction	11
2.2 NMX Server Authentication Configuration In BCM	11
2.2.1 NMX Manager Metrics Forwarding Via BCM And prometheusmetricforwarders	12
3 Rack Management And Management Policies	13
3.1 Introduction	13
3.2 Basic Rack Management	13
3.3 Introduction To Advanced Features In Rack Management	16
3.3.1 How BMS And BCM Work Together	16
3.3.2 Commands For Monitoring Advanced Features In Rack Management	18
3.3.3 Managing Leaks With rules In leakactionpolicies	20
4 BCM Power Shelf Integration	25
4.1 Power Shelf Listing And Overview	25
4.2 Power Shelves Networking Configuration	26
4.3 Access Configuration For The Power Management Controller	26
4.4 Power Shelf Settings Overview	27
4.5 Power Shelf Metrics Overview	27
4.6 Power Shelf Firmware	29
4.7 Power Shelf Circuit Monitoring	29
4.8 Power Shelves And Rack Mode Commands	29
4.8.1 The list Command	29
4.8.2 The rackoverview Command	30
4.8.3 The display Command	31
4.8.4 Power Management Of Racks	32
5 NVIDIA Autonomous Job Recovery	35
5.1 Introduction	35

5.2 Prerequisites To Install Autonomous Job Recovery: Kubernetes Installation	35
5.2.1 Kubernetes Components Needed For Installation Of Autonomous Job Recovery	35
5.2.2 Settings For Kubernetes Components Used In Autonomous Job Recovery In- stallation	36
5.3 Prerequisites To Install Autonomous Job Recovery: Slurm Installation	37
5.4 Autonomous Job Recovery Installation With <code>cm-mission-control-setup</code>	37
5.4.1 Credentials Required To Install Autonomous Job Recovery When Running <code>cm- mission-control-setup</code>	37
5.4.2 Running <code>cm-mission-control-setup</code>	37
5.5 Post-installation Checks	38
5.5.1 Grafana Access	39
5.5.2 Autonomous Job Recovery Failure When The Slurm Daemon Spool Directory Is Not Cleaned Up	39
6 NVIDIA Autonomous Hardware Recovery	41
6.1 Introduction	41
6.2 Prerequisites To Install Autonomous Hardware Recovery	41
6.3 Installation Of NVIDIA Autonomous Hardware Recovery Using <code>cm-mission- control-setup</code>	42
6.4 Verifying Autonomous Hardware Recovery Is Up And Ready	46
6.4.1 Checking The Pods	46
6.4.2 Testing The Web Interface	46
7 NVLink Switch Integration	51
7.1 NVLink Switch Listing And Overview	51
7.2 NVLink Switch Configuration	52
7.3 NVLink Switch Monitoring	54
7.4 NVLink Switch Redfish Events	54
8 Power Reservation Steering	55
8.1 Introduction	55
8.2 Deployment Of PRS	55
8.3 Changes Made On Deployment, And Managing Changes Post-Deployment With <code>cmsh</code>	57
8.3.1 PRS Values	57
8.3.2 Slurm Values	58
8.3.3 PRS Removal	59
8.4 Changes Made On Deployment, And Managing Changes Post-Deployment With Base View	59
8.4.1 Base View Power Reservation Steering Wizard	59
8.4.2 Base View Power Reservation Steering Client And Server Configuration Overlays	60
9 NVIDIA Mission Control DGX GB200 Measurables	61
9.1 Circuit-Related DGX GB200 Metrics	61
9.2 Leak Detection DGX GB200 Measurables	62
9.3 DGX GB200 NVLink Measurables	63
9.4 DGX GB200 Power Shelf Measurables	64
9.5 Cooling Distribution Unit Metrics On The DGX GB200	65
9.6 GPU Measurables For The DGX GB200	66

9.7 Prometheus Metrics For The DGX GB200	68
9.8 Redfish Metrics For The DGX GB200	70
9.9 Redfish Enums For The DGX GB200	86
9.10 Health Checks For The DGX GB200	91

Preface

Welcome to the *NVIDIA Mission Control Manual* for NVIDIA Base Command Manager 11.

0.1 About This Manual

The NVIDIA Mission Control features of NVIDIA Base Command Manager (BCM) are available on clusters that are running the NVIDIA Mission Control edition of BCM. The features are listed in chapter 1 of this manual.

This manual is aimed at helping cluster administrators install, understand, configure, and manage the NVIDIA Mission Control features and capabilities of NVIDIA Base Command Manager. The administrator is expected to be reasonably familiar with the *BCM Administrator Manual*. The *NVIDIA Mission Control Manual* is thus a supplementary manual to the *BCM Administrator Manual* for the NVIDIA Mission Control features.

0.2 About The Manuals In General

Regularly updated versions of the manuals are available on updated clusters by default at `/cm/shared/docs/cm`. The latest updates are always online at <https://docs.nvidia.com/base-command-manager>.

- The *Installation Manual* describes installation procedures for the basic cluster.
- The *Administrator Manual* describes the general management of the cluster.
- The *User Manual* describes the user environment and how to submit jobs for the end user.
- The *Cloudbursting Manual* describes how to deploy the cloud capabilities of the cluster.
- The *Developer Manual* has useful information for developers who would like to program with BCM.
- The *Edge Manual* explains how BCM can be used with edge sites.
- The *Containerization Manual* describes how to manage containers with BCM.

If the manuals are downloaded and kept in one local directory, then in most pdf viewers, clicking on a cross-reference in one manual that refers to a section in another manual opens and displays that section in the second manual. Navigating back and forth between documents is usually possible with keystrokes or mouse clicks.

For example: `<Alt>-<Backarrow>` in Acrobat Reader, or clicking on the bottom leftmost navigation button of `xpdf`, both navigate back to the previous document.

The manuals constantly evolve to keep up with the development of the BCM environment and the addition of new hardware and/or applications. The manuals also regularly incorporate feedback from administrators and users, who can submit comments, suggestions or corrections via the website

<https://enterprise-support.nvidia.com/s/create-case>

Section 14.2 of the *Administrator Manual* has more details on submitting an issue.

0.3 Getting Administrator-Level Support

If the reseller from whom BCM was bought offers direct support, then the reseller should be contacted.

Otherwise the primary means of support is via the website

<https://enterprise-support.nvidia.com/s/create-case>

Section 14.2 of the *Administrator Manual* has more details on working with support.

0.4 Getting Professional Services

The BCM support team normally differentiates between

- regular support (customer has a question or problem that requires an answer or resolution), and
- professional services (customer asks for the team to do something or asks the team to provide some service).

Professional services can be provided via the NVIDIA Enterprise Services page at:

<https://www.nvidia.com/en-us/support/enterprise/services/>

1 Introduction

1.1 NVIDIA Mission Control Requirements And Features

The NVIDIA Mission Control features are extended features that are only available in the NVIDIA Mission Control edition of NVIDIA Base Command Manager (BCM). NVIDIA Mission Control features are targeted at NVIDIA B200, GB200, and other OEM platforms. The NVIDIA sales team must be contacted in case an NVIDIA Mission Control edition is required.

The NVIDIA Mission Control features of the NVIDIA Mission Control edition of BCM can be enabled on a BCM cluster with an active BCM license. Activation of the BCM license is described in Chapter 4 of the *Installation Manual*.

1.1.1 NVIDIA Mission Control Features

The NVIDIA Mission Control edition is capable of running the following:

- NVIDIA-conforming BMS: a *Building Management System* software that aids in cluster management in the building.
- Leak detection: monitoring for leak detection in the fluid-cooled processors. Associated with this are the possible automated actions taken.
- Autonomous hardware recovery: Automates some hardware management operations in order to keep the cluster up.
- Autonomous job recovery: integrates a suite of autonomous software services to enhance job submission, restart procedures, carry out telemetry collection, and carry out anomaly resolution processes.
- NMX: Product suite name for HPC cluster network monitoring and management system that comprises NMX Telemetry, NMX Manager, NMX Controller and NMX Oasis.
- IMEX (Internode Memory Exchange Service): a secure service that facilitates the mapping of GPU memory over NVLink between the GPUs in an NVLink domain. IMEX configuration is described in section 7.5.1 of the *Administrator Manual* as part of workload management configuration.
- GB200 rack management: Automated management of rack power up and power down sequence.
- GB200 firmware management: GB200 firmware and BIOS management.
- Power reservation steering: Predicting power consumption, and capping the power to active components based on that prediction, to reduce data center energy costs.
- Run:ai: Kubernetes-based workload management and orchestration via Run:ai.

1.1.2 Checking That The NVIDIA Mission Control Edition Is Running

Whether the NVIDIA Mission Control edition is running on the cluster can be checked:

- in cmsh:
 - using the `licenseinfo` command within cmsh

Example

```
root@basecm11:~# cmsh -c 'main licenseinfo | egrep Version\|Edition'
Version                      10 and above
Edition                      NVIDIA Mission Control
```

or

- using the `verify-license` utility from a command line:

Example

```
root@basecm11:~# verify-license info
===== Certificate Information =====
Version:          10
Edition:          NVIDIA Mission Control
...
```

or

- using the `editiondisabledfeatures` command.

- * If the NVIDIA Mission Control edition is running, then the features are all available, as indicated by:

Example

```
root@basecm11:~# cmsh -c 'main editiondisabledfeatures'
No disabled features
```

- * If the NVIDIA Mission Control edition is not running, then the features are not available, as indicated by:

```
root@basecm11:~# cmsh -c 'main editiondisabledfeatures'
LeakDetection
DisplayRack
RackPower
AutonomousHardwareRecovery
AutonomousJobRecovery
Firmware.flash:GB200
BMS
Run:ai
IMEX
NMX
PRS
```

- in Base View the status of the NVIDIA Mission Control features can be seen in the navigation paths:

`Cluster > License Information > Mission Control`

or

`Cluster > Overview > License Information`

2 NMX Settings For NVLink Monitoring

2.1 Introduction

NMX is a product suite for HPC cluster network monitoring and management. The system is made up of:

- NMX Telemetry
- NMX Manager
- NMX Controller
- NMX Oasis

The NVIDIA Mission Control edition of BCM can monitor NMX telemetry services that run on:

- NVLinks (direct GPU interconnection within a server)
- NVLink switches (direct GPU interconnection across servers)

2.2 NMX Server Authentication Configuration In BCM

NMX service authentication settings can be configured from within `partition` mode, within the `nvmsettings` submode:

```
[basecm11->partition[base]->nmxmsettings]% show
Parameter          Value
-----
Revision
Server           10.140.0.41
User name        rw-user
Password         *****
Port             443
Verify SSL       no
CA certificate
Certificate
Private key
Prometheus metric forwarders   <0 in submode>
```

The corresponding Base View navigation path for this is using:

Cluster > Settings > BMC Settings > NMX settings > +

After BCM has provisioned the NMX Controller and the NMX Telemetry services on the NVSwitches, then the `sample_nmxm` data producer produces the NMX measurables. These can be viewed from `partition` mode:

Example

Measurable	Parameter	Type	Value	Age	State	Info
nmxm_chassis_count		NMX-M	9	30.9s		
nmxm_compute_allocation_count	all	NMX-M	54	30.9s		
nmxm_compute_allocation_count	free	NMX-M	48	30.9s		
nmxm_compute_allocation_count	full	NMX-M	6	30.9s		
nmxm_compute_allocation_count	partial	NMX-M	0	30.9s		
nmxm_compute_health_count	degraded	NMX-M	0	30.9s		
nmxm_compute_health_count	healthy	NMX-M	0	30.9s		
nmxm_compute_health_count	unhealthy	NMX-M	6	30.9s		
nmxm_compute_health_count	unknown	NMX-M	63	30.9s		
nmxm_compute_nodes_count		NMX-M	69	30.9s		
nmxm_domain_health_count	degraded	NMX-M	0	30.9s		
nmxm_domain_health_count	healthy	NMX-M	0	30.9s		
nmxm_domain_health_count	unhealthy	NMX-M	0	30.9s		
nmxm_domain_health_count	unknown	NMX-M	3	30.9s		
nmxm_gpu_health_count	degraded	NMX-M	0	30.9s		
nmxm_gpu_health_count	degraded_bw	NMX-M	0	30.9s		
nmxm_gpu_health_count	healthy	NMX-M	0	30.9s		
nmxm_gpu_health_count	nonvlink	NMX-M	24	30.9s		
nmxm_gpu_health_count	unknown	NMX-M	16	30.9s		
nmxm_gpus_count		NMX-M	40	30.9s		
nmxm_ports_count		NMX-M	152	30.9s		
nmxm_switch_health_count	healthy	NMX-M	0	30.9s		
nmxm_switch_health_count	missing_nvlink	NMX-M	8	30.9s		
nmxm_switch_health_count	unhealthy	NMX-M	0	30.9s		
nmxm_switch_health_count	unknown	NMX-M	0	0.9s		
nmxm_switch_nodes_count		NMX-M	46	30.9s		

2.2.1 NMX Manager Metrics Forwarding Via BCM And prometheusmetricforwarders

Prometheus servers can be used to pick up NMX Manager metrics via Prometheus exporters. However in a typical BCM configuration, such as in a BCM type 1 network (section 3.3.9 of the *Installation Manual*), the Prometheus exporter is behind the BCM head node firewall. This means that by default an external Prometheus server cannot reach the exporter endpoint.

In such a case the `prometheusmetricforwarders` submode can be used to define the endpoint URL and HTTP method that the exporter uses. This forwards the exporter endpoint to the external network interface of the head node. The external Prometheus server can then reach the forwarded exporter endpoint.

Example

Parameter	Value
Name	all
Revision	
Urls	http://10.140.0.41:9091/metrics
Method	GET

The corresponding Base View navigation path for configuring Prometheus metrics forwarding is using:

Cluster > Settings > BMC Settings > NMX settings >  > Prometheus metric forwarders
> 

3 Rack Management And Management Policies

3.1 Introduction

Rack management is the process of using tools and software to organize and manage the infrastructure components and racks in a data center.

A DGX GB200 system can be thought of as a “disaggregated supercomputer in a rack”. The cluster administrator can manage the individual nodes in the traditional way, but BCM rack management features introduced in BCM version 11 allow the GB200 nodes to be managed more conveniently as unified NVL domains.

3.2 Basic Rack Management

Basic rack management is covered in section 3.16 of the *Administrator Manual*. The commands in **rack** mode that are discussed there include:

- The `list` command, which can give useful information on how the devices in racks are organized, if the attributes of racks in **rack** mode have been set up in a meaningful way. An example is given in section 4.8.1.
- The `display` command, which is useful for seeing the position of where devices are located in the rack, if the cluster administrator has recorded the device positions in the rack. An example is given in section 4.8.3.
- The `addrackposition` command (page 184 of the *Administrator Manual*) was introduced in BCM version 11. It allows rack positions in a cluster to be filled more easily than with a looping script.
- The `rackoverview` command (page 180 of the *Administrator Manual*) was also introduced in BCM version 11. It allows rack positions in a cluster to be overviewed conveniently. It shows information about the types of entities in a rack. It also then lists some more detailed information about some of the entity types.

For a DGX GB200 platform that has been configured, the output of the `rackoverview` command looks similar to the following for a specified rack (here it is a rack named b05 that has 18 nodes, 9 switches, and 8 power shelves):

Example

Type	Up	Down	Closed	Total
Nodes	18	0	0	18

DPU nodes	0	0	0	0		
Managed switches	0	0	0	0		
NVLink switches	9	0	0	9		
Power shelves	8	0	0	8		
Devices	0	0	0	0		
Cores	2,592	-	-	2,592		
GPUs	72	-	-	72		
<hr/>						
Name	Value					
<hr/>						
User CPU	0.03%					
System CPU	0.07%					
Idle CPU	99.9%					
Other CPU	0.0%					
Memory used	1.09 TiB (3.75%)					
Memory unused	28.4 TiB (97.9%)					
Memory total	29.0 TiB					
Total GPU utilization	0 W					
Total GPU power usage	11.6 KW					
Total GPU NVlink bandwidth	0 W					
Average GPU temperature	30.9028 C					
Total CPU power usage	1.60 KW					
Average CPU temperature	48.6361 C					
<hr/>						
Node	GPU	Utilization	Temperature	Power usage	Memory used...	
<hr/>					...	
b05-p1-dgx-05-c01	gpu0	0.0%	30 C	167.184 W	0 B	...
b05-p1-dgx-05-c01	gpu1	0.0%	31 C	154.726 W	0 B	...
b05-p1-dgx-05-c01	gpu2	0.0%	32 C	154.894 W	0 B	...
b05-p1-dgx-05-c01	gpu3	0.0%	31 C	170.649 W	0 B	...
b05-p1-dgx-05-c02	gpu0	0.0%	31 C	162.994 W	0 B	...
b05-p1-dgx-05-c02	gpu1	0.0%	31 C	151.328 W	0 B	...
b05-p1-dgx-05-c02	gpu2	0.0%	31 C	151.149 W	0 B	...
b05-p1-dgx-05-c02	gpu3	0.0%	31 C	170.537 W	0 B	...
b05-p1-dgx-05-c03	gpu0	0.0%	31 C	166.664 W	0 B	...
b05-p1-dgx-05-c03	gpu1	0.0%	31 C	154.725 W	0 B	...
b05-p1-dgx-05-c03	gpu2	0.0%	31 C	151.047 W	0 B	...
b05-p1-dgx-05-c03	gpu3	0.0%	31 C	170.125 W	0 B	...
b05-p1-dgx-05-c04	gpu0	0.0%	31 C	170.642 W	0 B	...
b05-p1-dgx-05-c04	gpu1	0.0%	31 C	155.034 W	0 B	...
b05-p1-dgx-05-c04	gpu2	0.0%	31 C	152.078 W	0 B	...
b05-p1-dgx-05-c04	gpu3	0.0%	31 C	166.551 W	0 B	...
b05-p1-dgx-05-c05	gpu0	0.0%	31 C	170.516 W	0 B	...
b05-p1-dgx-05-c05	gpu1	0.0%	31 C	158.7 W	0 B	...
b05-p1-dgx-05-c05	gpu2	0.0%	31 C	158.7 W	0 B	...
b05-p1-dgx-05-c05	gpu3	0.0%	31 C	174.413 W	0 B	...
b05-p1-dgx-05-c06	gpu0	0.0%	30 C	166.44 W	0 B	...
b05-p1-dgx-05-c06	gpu1	0.0%	30 C	154.628 W	0 B	...
b05-p1-dgx-05-c06	gpu2	0.0%	31 C	154.865 W	0 B	...
b05-p1-dgx-05-c06	gpu3	0.0%	31 C	170.423 W	0 B	...
b05-p1-dgx-05-c07	gpu0	0.0%	31 C	164.078 W	0 B	...
b05-p1-dgx-05-c07	gpu1	0.0%	31 C	158.303 W	0 B	...
b05-p1-dgx-05-c07	gpu2	0.0%	31 C	150.958 W	0 B	...
b05-p1-dgx-05-c07	gpu3	0.0%	31 C	170.537 W	0 B	...
b05-p1-dgx-05-c08	gpu0	0.0%	31 C	166.661 W	0 B	...
b05-p1-dgx-05-c08	gpu1	0.0%	31 C	155.875 W	0 B	...
b05-p1-dgx-05-c08	gpu2	0.0%	31 C	154.726 W	0 B	...
b05-p1-dgx-05-c08	gpu3	0.0%	32 C	170.311 W	0 B	...
b05-p1-dgx-05-c09	gpu0	0.0%	31 C	166.331 W	0 B	...

b05-p1-dgx-05-c09	gpu1	0.0%	31 C	155.469 W	0 B	...
b05-p1-dgx-05-c09	gpu2	0.0%	30 C	154.751 W	0 B	...
b05-p1-dgx-05-c09	gpu3	0.0%	32 C	175.429 W	0 B	...
b05-p1-dgx-05-c10	gpu0	0.0%	30 C	174.181 W	0 B	...
b05-p1-dgx-05-c10	gpu1	0.0%	31 C	158.523 W	0 B	...
b05-p1-dgx-05-c10	gpu2	0.0%	31 C	152.808 W	0 B	...
b05-p1-dgx-05-c10	gpu3	0.0%	31 C	166.551 W	0 B	...
b05-p1-dgx-05-c11	gpu0	0.0%	30 C	166.704 W	0 B	...
b05-p1-dgx-05-c11	gpu1	0.0%	31 C	154.726 W	0 B	...
b05-p1-dgx-05-c11	gpu2	0.0%	31 C	155.034 W	0 B	...
b05-p1-dgx-05-c11	gpu3	0.0%	30 C	170.679 W	0 B	...
b05-p1-dgx-05-c12	gpu0	0.0%	31 C	170.424 W	0 B	...
b05-p1-dgx-05-c12	gpu1	0.0%	31 C	155.601 W	0 B	...
b05-p1-dgx-05-c12	gpu2	0.0%	31 C	150.892 W	0 B	...
b05-p1-dgx-05-c12	gpu3	0.0%	31 C	166.806 W	0 B	...
b05-p1-dgx-05-c13	gpu0	0.0%	31 C	166.766 W	0 B	...
b05-p1-dgx-05-c13	gpu1	0.0%	30 C	151.894 W	0 B	...
b05-p1-dgx-05-c13	gpu2	0.0%	31 C	151.22 W	0 B	...
b05-p1-dgx-05-c13	gpu3	0.0%	31 C	170.205 W	0 B	...
b05-p1-dgx-05-c14	gpu0	0.0%	31 C	172.778 W	0 B	...
b05-p1-dgx-05-c14	gpu1	0.0%	31 C	162.152 W	0 B	...
b05-p1-dgx-05-c14	gpu2	0.0%	31 C	150.857 W	0 B	...
b05-p1-dgx-05-c14	gpu3	0.0%	31 C	166.33 W	0 B	...
b05-p1-dgx-05-c15	gpu0	0.0%	31 C	166.518 W	0 B	...
b05-p1-dgx-05-c15	gpu1	0.0%	31 C	154.729 W	0 B	...
b05-p1-dgx-05-c15	gpu2	0.0%	31 C	154.1 W	0 B	...
b05-p1-dgx-05-c15	gpu3	0.0%	31 C	170.537 W	0 B	...
b05-p1-dgx-05-c16	gpu0	0.0%	30 C	174.067 W	0 B	...
b05-p1-dgx-05-c16	gpu1	0.0%	30 C	158.385 W	0 B	...
b05-p1-dgx-05-c16	gpu2	0.0%	31 C	151.157 W	0 B	...
b05-p1-dgx-05-c16	gpu3	0.0%	31 C	166.77 W	0 B	...
b05-p1-dgx-05-c17	gpu0	0.0%	31 C	166.161 W	0 B	...
b05-p1-dgx-05-c17	gpu1	0.0%	31 C	158.62 W	0 B	...
b05-p1-dgx-05-c17	gpu2	0.0%	31 C	151.058 W	0 B	...
b05-p1-dgx-05-c17	gpu3	0.0%	31 C	169.704 W	0 B	...
b05-p1-dgx-05-c18	gpu0	0.0%	31 C	166.441 W	0 B	...
b05-p1-dgx-05-c18	gpu1	0.0%	31 C	155.164 W	0 B	...
b05-p1-dgx-05-c18	gpu2	0.0%	31 C	150.758 W	0 B	...
b05-p1-dgx-05-c18	gpu3	0.0%	31 C	169.752 W	0 B	...
<hr/>						
Switch	Utilization	Temperature	Power usage	Fan speed	Links active	L...
<hr/>						
B05-P1-NVSW-01	30.2%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-02	27.5%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-03	0.0%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-04	27.6%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-05	27.6%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-06	27.5%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-07	27.5%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-08	27.5%	0 C	0 W	0 RPM	0	0
B05-P1-NVSW-09	27.6%	0 C	0 W	0 RPM	0	0
<hr/>						
Power shelf	Input power	Output power	Temperature	Fan speed	Active PSU	To...
<hr/>						
B05-P1-PWR-01	3.7 KW	3.5 KW	43.3333 C	6.7 KRPM	6	6
B05-P1-PWR-02	3.11 KW	2.95 KW	42.4 C	6.7 KRPM	6	6
B05-P1-PWR-03	3.7 KW	3.5 KW	44.1667 C	6.7 KRPM	6	6
B05-P1-PWR-04	3.8 KW	3.6 KW	44 C	6.7 KRPM	6	6
B05-P1-PWR-05	3.6 KW	3.4 KW	45.5 C	6.7 KRPM	6	6

B05-P1-PWR-06	3.8 KW	3.6 KW	47.1667 C	6.7 KRPM	6
B05-P1-PWR-07	3.5 KW	3.4 KW	46 C	6.7 KRPM	6
B05-P1-PWR-08	3.6 KW	3.5 KW	47 C	6.7 KRPM	6

[a03-p1-head-01->rack]%

3.3 Introduction To Advanced Features In Rack Management

Over time, clusters have become more powerful and their cooling requirements have become more demanding. Liquid cooling technology has therefore become more common over time.

The monitoring associated with liquid cooling has also developed. Monitoring of the electrical supply, fans, component temperatures, and coolant values such as flow and pressure has become popular. Coolant monitoring in BCM includes an ability to detect leaks and then take actions. An NVIDIA-conforming BMS (section 3.3.1) provides some of these monitoring and managing features, including features associated with liquid cooling and electrical isolation.

The cluster administrator can monitor these advanced features using mostly rack mode commands (section 3.3.2).

3.3.1 How BMS And BCM Work Together

A Building Management System (BMS) is a software used to manage IT assets in a building. A BMS software that can integrate with BCM, is known as a *conforming BMS*. The conformance is to a reference design for data center planning (NVIDIA Controls and Monitoring Reference Design), available to some NVIDIA product users.

Since BCM version 11, BCM integrates with a conforming BMS for leak detection and control in the data center (DC). The *NVIDIA Mission Control Software Installation Guide* describes the implementation of the integration between BCM and a conforming BMS at:

<https://docs.nvidia.com/mission-control/docs/nmc-software-installation-guide/2.0.0/integration-of-bms-with-bcm.html>

The guide is also meant to serve as a guide to NVIDIA customers who may use different BMS solutions.

BCM communicates with a conforming BMS. Together they monitor the DC. The integration works with DGX GB200 systems that are liquid-cooled, as well as other systems.

In a DC, a conforming BMS is able to detect most of the detected leak-related events. The remaining detected leak-related events, from the in-tray BMCs (compute tray and switch tray), are dealt with by BCM.

BCM uses the MQTT protocol for two-way communication with the conforming BMS, using a BCM service `cm-mqtt`.

The BMS sends BCM information about the inventory of the DC, and about the leak-related events it knows about (non-tray events). MQTT topics are created, and structured in a way that BCM can parse.

BCM in turn sends the BMS instructions on liquid management or power management, according to defined policies on events such as leaks or power outages (section 4.7).

Building Management System Configuration

BMS integration can be configured in BCM in partition mode in cmsh and Base View.

The type of BMS can be configured in partition mode:

Example

```
root@basecm11:~# cmsh
[basecm11]% partition
[basecm11->partition[base]]% set bms nvidia-conforming\ bms; commit
```

In the preceding example, the BMS is set to an NVIDIA-conforming BMS. The possible types of BMS configuration can be:

- NVIDIA-conforming BMS
- pipe
- file
- URL

If bms is set to the URL of a BMS server, for example `http://bmserver/path`, then credentials for it can then be set as follows:

Example

```
root@basecm11:~# cmsh
[basecm11]# partition
[basecm11->partition[base]]% set bms url
[basecm11->partition*[base*]]% set bmssertificate <BMS certificate>
[basecm11->partition*[base*]]% set bmspath http://bmsserver/path
[basecm11->partition*[base*]]% set bmsprivatekey <BMS private key>
[basecm11->partition*[base*]]% commit
```

In the following example, an mqtt role is assigned in the mqtt configuration overlay. The overlay specifies where the cm-mqtt service should run. In this case it runs on all the head nodes:

Example

```
root@basecm11:~# cmsh
[basecm11->partition[base]]% configurationoverlay
[basecm11->configurationoverlay]% add mqtt
[basecm11->configurationoverlay*[mqtt*]]% set allheadnodes yes
```

Within the role assignment, the administrator specifies the hostname or IP address of where the BMS server runs, and where the cm-mqtt service pulls data from:

```
[basecm11->configurationoverlay*[mqtt*]]% roles
[basecm11->configurationoverlay*[mqtt*]->roles]% assign mqtt
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]]% show
Parameter          Value
-----
Name              mqtt
Revision
Type              MQTTRole
Add services      yes
Servers           <0 in submode>
CA certificate path /cm/local/apps/cmd/pythoncm/lib/python3.12/site-packages/\
pythoncm/etc/cacert.pem
Private key path  /cm/local/apps/cmd/cm-mqtt/etc/mqtt.key
Certificate path  /cm/local/apps/cmd/cm-mqtt/etc/mqtt.pem
Write named pipe path /var/spool/cmd/mqtt.pipe
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]]% servers
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]->servers]% list
Server (key)      Port      Disabled
----- -----
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]->servers]% add mqhost
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]->servers*[mqhost*]]% show
Parameter          Value
-----
```

```

Revision
Server          mqhost
Port            1883
Topic           BCM/#
Disabled        no
Username
Password       < not set >
Transport      tcp
Protocol       v3.1.1
Certificate required yes
CA certificate
Certificate
Private key
[basecm11->configurationoverlay*[mqtt*]->roles*[mqtt*]->servers*[mqhost*]]% set certificate
...

```

A username, password, and service certificate can be set. These allow BCM to communicate with the BMS securely:

```

[basecm11->configurationoverlay[mqtt]->roles[mqtt]->servers[mqhost]]% show
Parameter          Value
-----
Revision
Server          mqhost
Port            1883
Topic           BCM/#
Disabled        no
Username        cronus
Password        *****
Transport      tcp
Protocol       v3.1.1
Certificate required yes
CA certificate
Certificate
Private key
[basecm11->configurationoverlay[mqtt]->roles[mqtt]->servers[mqhost]]%

```

If the cluster manager has correctly established communication with the BMS, then the status of the BMS (metrics for power circuit, CDUs, electrical, and liquid cooling) can be seen from within cmsh (section 4.7).

Example

```

[basecm11]% powercircuit overview
Power circuit  Power    Current   Current limit   Current phase 1   Current phase 2 ...
-----
RPP-B12-3     14.0 KW  26.1118 A  0 A           17.7822 A      26.1118 A      ...
RPP-B14-3     15.3 KW  29.6178 A  0 A           17.7419 A      29.6178 A      ...
RPP-B21-5     7.5 KW   11.2255 A  0 A           10.4872 A      10.5412 A      ...
...

```

3.3.2 Commands For Monitoring Advanced Features In Rack Management

The electricaloverview Command

The `electricaloverview` command within rack mode gives a convenient overview of the electrical supply to the cluster.

The power used and the power budget of a specified rack (B05 in the following example) can be seen:

Example

```
[basecm11->rack[B05]]% electricaloverview
Rack          Power used      Power budget     Isolation status
-----
B05           28.5 kW        135 kW          0
```

If no rack is specified, then the command displays the information for all racks.

The liquidcoolingoverview Command

The `liquidcoolingoverview` command can be run from within `rack` mode, and conveniently displays coolant metrics:

```
[basecm11->rack]% liquidcoolingoverview a05
Supply          Return         Flow          Differential   Isolation
Rack  temperature  temperature  rate          pressure       status
-----
A05   26.03 C     30.78 C     70.938 LPM  31 KPa        0
```

If no rack is specified, then the command displays the information for all racks.

The leakdetectionoverview Command

The `leakdetectionoverview` command within `rack` mode gives a convenient overview about sensor readings related to leak detection in BCM and the electrical isolation status.

The output from the command applied to a specific rack looks similar to the following:

```
[basecm11->rack]% leakdetectionoverview a05
                           Liquid          Liquid          Electrical  Electrical
                           Leak           Leak           isolation    Liquid    isolation  isolation  Electrical
                           tray          sensor request   (BCM)      (BMS)    status     (BCM)    (BMS)    status
Rack detected detected fault
-----
A05   0            0            0            0            0            0            0            0            0
[basecm11->rack]%
```

If no rack is specified, then the command displays the information for all racks.

The leakinfo Command

The `leakinfo` command, is related to leakage detection, but it is not a `rack` mode command. It is a `device` mode command. It is noted in this section for completeness. It can be run at `device` mode level to list all leaks detected in the cluster.

Typically, and ideally, running the command shows something similar to:

```
[basecm11->device]% leakinfo
Hostname      Timestamp Severity  Info
-----
[basecm11->device]%
```

If there are issues on the devices, then the output may look similar to:

```
[basecm11->device]% leakinfo
Hostname      Timestamp Severity  Info
-----
a05-dgx-01-c09 15:09:23 large    The state of resource `Chassis_0_LeakDete...
a07-dgx-03-c16 15:38:48 large    The state of resource `Chassis_0_LeakDete...
b05-dgx-05-c08 09:11:06 large    The state of resource `Chassis_0_LeakDete...
b05-dgx-05-c10 09:11:05 large    The state of resource `Chassis_0_LeakDete...
```

```
b06-dgx-06-c17 11:49:30 large      The state of resource `Chassis_0_LeakDete...
[basecm11->device]%
```

Running `help leakinfo` lists further command options.

The isolationrequestinfo Command

If a leak is detected, then BCM receives the report. BCM then:

- powers off the node that is leaking
- powers off the rack if multiple nodes are leaking
- isolates the rack to prevent other nodes suffering the same fate

Isolation means that power no longer runs through the leaking racks.

The isolation status of a rack can be viewed with the `isolationrequestinfo` command, which normally shows something similar to:

Example

```
[basecm11->rack]% isolationrequestinfo
Rack          Timestamp   Kind     Value   Info
-----
[basecm11->rack]%
```

During a leakage event, the output may show something similar to:

Example

```
[basecm11->rack]% isolationrequestinfo
Rack  Timestamp Kind Value Info
-----
B05  03:40:01    4    Default building policy/Rack
[basecm11->rack]%
```

3.3.3 Managing Leaks With rules In leakactionpolicies

The cluster administrator has the freedom to construct actions based on triggers based on leak detection measurables. While that is powerful, that can be quite labor-intensive for cluster administrators.

BCM provides a simpler way to deal with leaks, based on leak action policies. Leak action policies are a collection of rules that carry out actions to be followed with leaks are detected. Racks play an important part in these policies and rules.

The policies use predefined scopes, or groupings. The names of these groupings are rules. The groupings are listed next, in order from largest to smallest, along with their associated descriptions:

- Building: a building is made up of rooms
- Room: A room has rows of racks in it.
- Row: A row of racks has individual racks in it.
- Rack: An individual rack has devices in it.
- Device: The smallest possible grouping

The policies can be listed in the `leakactionpolicies` submode, under partition mode.

Example

```
[basecm11->partition[base]->leakactionpolicies]% list
Name (key)          Rules
-----
Default building policy Building, Device, Rack, Room, Row
Default device policy   Device
Default rack policy     Device, Rack
Default room policy     Device, Rack, Room, Row
Default row policy      Device, Rack, Row
```

The rules used by each policy are also shown, under the Rules column in the preceding example.

Each rule in a policy is a copy of a template. By default the rules per policy are identical.

Just the rules used by a specific policy can also be viewed:

Example

```
[basecm11->partition[base]->leakactionpolicies]% use default building policy
[basecm11->partition[base]->leakactionpolicies[Default building policy]]% show
Parameter          Value
-----
Name               Default building policy
Revision
Rules              Building, Device, Rack, Room, Row
```

Within a policy, the value of the settings for these rules can be listed:

Example

```
[basecm11->...leakactionpolicies[Default building policy]->rules]% list
                           Power  Power  Electrical  Liquid  Minimal  Maximal
Name (key)  Scope    off    off  rack isolation isolation devices devices
-----
Building    building yes    yes    no        no       11      100000
Device      device   yes    no     no        no
Rack        rack    yes    yes    yes        yes      3       100
Room        room    yes    yes    no        no       7      100000
Row         row     yes    yes    no        no       5       100
```

The rules for each policy can be modified independently for each policy, so that, for example the Rack rule values may differ between Default row policy and Default building policy. However, a sensible cluster administrator is not expected to carry out such a modification.

Another way of phrasing this is: The values of the rules can be changed from the defaults. In that case, a sensible administrator is expected to keep the values synchronized across the rules in the leak action policies.

In the preceding example:

If the minimal devices criterion is met for leaking for each rule in a policy, then the values for the column header are applied to the rules of that policy.

So, for the policy rules in the preceding example, when leaking is detected:

- If only one or two devices are detected as leaking, then only the devices involved are powered off.
- If, in one rack, more than 3 devices, but less than 5 devices, are detected as leaking, then in addition to the device rules being applied, the entire rack is powered off. Isolation is also carried out for the electrical power and liquid coolant.

- If more devices meet the conditions, so a larger scope is affected by leaking, then the rules for the larger scope are also applied.

The rules for a specific policy can be removed. The `row` rule for Default building policy can be removed with, for example:

```
[basecm11->partition[base]->leakactionpolicies[Default building policy]]% rules
[basecm11->...leakactionpolicies[Default building policy]->rules]% remove <TAB><TAB>
building device rack room row
[basecm11->...leakactionpolicies[Default building policy]->rules]% remove row; commit
```

Custom Rules For leakactionpolicies

The names for the rules (Device, Rack, Row ...) are really just default labels, and a sensible cluster administrator is expected to have the description associated with the name match the physical reality. Creating a custom rule may help with this.

Custom rules can be added by the cluster administrator, by first adding another policy:

```
[basecm11->partition[base]->leakactionpolicies[Default building policy]]% add mypolicy
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]]% list
Name (key)          Rules
-----
Default building policy Building, Device, Rack, Room, Row
Default device policy Device
Default rack policy   Device, Rack
Default room policy   Device, Rack, Room, Row
Default row policy    Device, Rack, Row
mypolicy            none
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]]% show
Parameter          Value
-----
Name                mypolicy
Revision           none
Rules              none
```

By default, there are no rules in the new policy:

```
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]]% rules
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]->rules]%
Name (key)  Scope      Power off  Power off  Electrical  Liquid  Minimal  Maximal
          rack      isolation  isolation  devices  devices
```

The cluster administrator may have a testing room in which to keep a rack. In that case, the following might be added:

```
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]->rules]%
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]->rules*[testroom*]]% show
Parameter          Value
-----
Name                testroom
Revision           rack
Scope               no
Disabled           1
Minimal devices    100
Maximal devices    0s
Grace period       5m
Correlation window yes
Power off          yes
```

Power off rack	no
Electrical isolation	no
Liquid isolation	no
Minimal severity	1
Maximal severity	100

Adding a rule adds a default with the preceding rack-based rule values.

The rule values can be modified using the set command.

Running set without any options provides a help text explaining what can be set:

```
[basecm11->partition*[base*]->leakactionpolicies*[mypolicy*]->rules]% set
```


4 BCM Power Shelf Integration

The NVIDIA GB200 server platforms use power shelves to distribute power in racks. Several power shelves are used per rack. Each power shelf takes up one rack unit. Each power shelf comprises several PSUs (Power Supply Units).

The reference design for an NVIDIA GB200 rack has 6 power shelves, and uses 6 PSUs in each power shelf. Each PSU is rated up to 5.5kW, so that each power shelf is rated up to 33kW.

This chapter considers an NVIDIA GB200 server platform that has been configured to work with BCM.

4.1 Power Shelf Listing And Overview

The `list -t powershelf` command in device mode can be used to list the power shelves of a configured cluster

Power shelves are configured as devices on an IPMI network, as indicated by the following example (some rows ellipsized):

Example

Type	Hostname (key)	MAC	IP	Network	Status
PowerShelf	B05-P1-PWR-01	00:18:23:0C:1E:90	7.241.1.47	rfnet	[UP]
...					
PowerShelf	B05-P1-PWR-08	00:18:23:0C:1E:85	7.241.1.54	rfnet	[UP]
PowerShelf	a05-p1-pwr-01	00:18:23:0C:1E:84	7.241.1.8	rfnet	[UP]
...					
PowerShelf	a05-p1-pwr-08	00:18:23:0C:1E:8F	7.241.1.80	rfnet	[UP]
PowerShelf	b06-p1-pwr-01	00:18:23:0C:40:2A	7.241.1.55	rfnet	[UP]
...					
PowerShelf	b06-p1-pwr-08	00:18:23:0C:40:99	7.241.1.62	rfnet	[UP]

The `powershelfoverview` command displays an overview of a particular power shelf, showing the measurements from its PSUs (6 in the following example):

Example

	Input Supply power	Input current	Input voltage	Rail power	Rail current	Rail voltage	Fan Temperature	Fan speed	Fan Healthy
1	679 W	2.98 A	240.75 V	652 W	13.06 A	49.85 V	44 C	6.7 KRPM	PASS

2	632	W	2.8	A	240.75	V	607	W	12.18	A	49.87	V	44	C	6.6	KRPM	PASS
3	599	W	2.64	A	240.75	V	575	W	11.48	A	49.87	V	43	C	6.8	KRPM	PASS
4	599	W	2.63	A	240.75	V	575	W	11.51	A	49.89	V	43	C	6.8	KRPM	PASS
5	617	W	2.76	A	240.5	V	574	W	11.5	A	49.87	V	43	C	6.9	KRPM	PASS
6	609	W	2.64	A	240.5	V	584	W	11.7	A	49.87	V	43	C	6.6	KRPM	PASS

4.2 Power Shelves Networking Configuration

The network interface settings for a specific power shelf can be managed with its Redfish interface (section 3.7.1 of the *Administrator Manual*). For example, a power shelf, here named B05-P1-PWR-01, can have the following values set with the `set` command:

Example

```
[basecm11->device[B05-P1-PWR-01]->interfaces[rf0]]% show
Parameter          Value
-----
Revision
Type              bmc
Network device name rf0
Network            rfnet
IP                7.241.1.47
DHCP               no
Alternative Hostname
Additional Hostnames
Switch ports
Start if           always
BringUpDuringInstall no
On network priority 60
Bootable            no
MAC                00:18:23:0C:1E:90
Speed
Card Type
```

Redfish metrics, which are the ones prefixed by RF_ (as seen in section 9.8), are sampled only when the BMC interface starts with rf0, rf1, ...

4.3 Access Configuration For The Power Management Controller

Access to the Power Management Controller (PMC) is configured in the BMC settings:

Example

```
[basecm11->device[B05-P1-PWR-01]->bmcsettings]]% show
Parameter          Value
-----
Revision
User name          root
Password           *****
User ID             -1
Power reset delay 0s
Extra arguments
Privilege          administrator
```

Authentication allows NVIDIA Mission Control to monitor the power shelves using Redfish sampling.

4.4 Power Shelf Settings Overview

An overview of the settings for a particular power shelf can be viewed with the `show` command:

Example

```
[basecm11->device[B05-P1-PWR-01]]% show
Parameter          Value
-----
Hostname           B05-P1-PWR-01
IP                7.241.1.47
Network            rfnet
Revision
Type              PowerShelf
Mac               00:18:23:0C:1E:90
Model
Activation        Tue, 11 Feb 2025 16:45:39 PST
Rack              B05:6
Chassis            < not set >
Access Settings   <submode>
Management network rfnet
Power control     none
Custom power script
Custom power script argument
Power distribution units
Default gateway metric      0
Switch ports          B04-P1-00B-02:1
Interfaces            <1 in submode>
PMC Settings          <submode>
Userdefined1
Userdefined2
User defined resources
Supports GNSS        no
Custom ping script
Custom ping script argument
Partition            base
Part number
Serial number
Notes               <0B>
Prometheus metric forwarders <0 in submode>
```

4.5 Power Shelf Metrics Overview

A listing of the power shelf metrics can be seen with some grepping (the output for PSUs 2 to 5 have been ellipsized for readability in this manual):

```
[basecm11->monitoring->measurable]% list -f name:40,parameter:9,producer
|head -2; list -f name:40,parameter:9,producer| egrep -i
'(powershel|redfish_power_shelf)'
name (key)          parameter producer
-----
AveragePowerShelfTemperature          AggregatePowerShelf
PowerShelvesClosed                  ClusterTotal
PowerShelvesDown                   ClusterTotal
PowerShelvesTotal                  ClusterTotal
PowerShelvesUp                     ClusterTotal
RF_Active_PSU                      redfish_power_shelf
RF_PowerShelf_Status               1             redfish_power_shelf
...                                ...
```

RF_PowerShelf_Status	6	redfish_power_shelf
RF_PowerShelf_Status	total	redfish_power_shelf
RF_Power_Supply_Energy	1	redfish_power_shelf
...		
RF_Power_Supply_Energy	6	redfish_power_shelf
RF_Power_Supply_Energy	total	redfish_power_shelf
RF_Power_Supply_FanSpeed	1	redfish_power_shelf
...		
RF_Power_Supply_FanSpeed	6	redfish_power_shelf
RF_Power_Supply_FanSpeed	average	redfish_power_shelf
RF_Power_Supply_InputCurrent	1	redfish_power_shelf
...		
RF_Power_Supply_InputCurrent	6	redfish_power_shelf
RF_Power_Supply_InputCurrent	total	redfish_power_shelf
RF_Power_Supply_InputPower	1	redfish_power_shelf
...		
RF_Power_Supply_InputPower	6	redfish_power_shelf
RF_Power_Supply_InputPower	total	redfish_power_shelf
RF_Power_Supply_InputVoltage	1	redfish_power_shelf
...		
RF_Power_Supply_InputVoltage	6	redfish_power_shelf
RF_Power_Supply_LifetimeReading_Energy	1	redfish_power_shelf
...		
RF_Power_Supply_LifetimeReading_Energy	6	redfish_power_shelf
RF_Power_Supply_OutputPower	1	redfish_power_shelf
...		
RF_Power_Supply_OutputPower	6	redfish_power_shelf
RF_Power_Supply_OutputPower	total	redfish_power_shelf
RF_Power_Supply_PowerFactor_Input	1	redfish_power_shelf
...		
RF_Power_Supply_PowerFactor_Input	6	redfish_power_shelf
RF_Power_Supply_RailCurrent	1	redfish_power_shelf
...		
RF_Power_Supply_RailCurrent	6	redfish_power_shelf
RF_Power_Supply_RailCurrent	total	redfish_power_shelf
RF_Power_Supply_RailPower	1	redfish_power_shelf
...		
RF_Power_Supply_RailPower	6	redfish_power_shelf
RF_Power_Supply_RailPower	total	redfish_power_shelf
RF_Power_Supply_RailVoltage	1	redfish_power_shelf
...		
RF_Power_Supply_RailVoltage	6	redfish_power_shelf
RF_Power_Supply_Temperature	1	redfish_power_shelf
...		
RF_Power_Supply_Temperature	6	redfish_power_shelf
RF_Power_Supply_Temperature	average	redfish_power_shelf
RF_Power_Supply_code_error	1	redfish_power_shelf
...		
RF_Power_Supply_code_error	6	redfish_power_shelf
RF_Power_Supply_message_error	1	redfish_power_shelf
...		
RF_Power_Supply_message_error	6	redfish_power_shelf
RF_Summary_Metrics_Power		redfish_power_shelf
RF_Summary_Metrics_Temperature		redfish_power_shelf
RF_Total_PSU		redfish_power_shelf
TotalPowerShelfInputCurrent		AggregatePowerShelf
TotalPowerShelfInputPower		AggregatePowerShelf
TotalPowerShelfRailCurrent		AggregatePowerShelf
TotalPowerShelfRailPower		AggregatePowerShelf

The descriptions for these metrics are as indicated in table 9.4.

4.6 Power Shelf Firmware

Power shelf firmware is managed by the `firmware` command options as described in Section 14.5.3 of the *Administrator Manual*.

4.7 Power Shelf Circuit Monitoring

A power circuit typically supplies several power shelves, based on the rated power of the circuit and power shelves. Administration and wiring convenience may also play a role in allocating a circuit to a power shelf.

So, for example, the administrator may configure a rack with 6 power shelves so that it is supplied power from one power circuit.

The power circuit breakers used by each power circuit, typically provided by a remote power panel (RPP), can have their locations noted for administrative convenience.

Example

```
[basecm11->powercircuit]# list
Name (key) Row Room Building Location Racks
-----
RPP-B12-3 1 G Watchtower Trantor
RPP-B14-3
RPP-B21-5
RPP-B21-6
RPP-B22-5
...
```

The cluster may have been set up to monitor the power circuit breaker for the power reaching BCM (section 3.3.1). In that case, the `overview` command typically displays the power consumption and current per

Example

```
[basecm11->powercircuit]# overview RPP-B14-3
phase.          Current Current Current Current
Power circuit Power Current limit phase 1 phase 2 phase 3 Racks
-----
RPP-B14-3     15.6 KW 30.137 A 0 A    18.28 A 30.119 A 28.134 A 1
```

A value of 0 A means that the current limit is not reported.

4.8 Power Shelves And Rack Mode Commands

This section discusses some of the commands that can be useful for setting up and viewing devices from within the `rack` mode of `cmsh` (section 3.16 of the *Administrator Manual*), but with a focus on power shelf management.

4.8.1 The `list` Command

If the attributes of racks in `rack` mode have been set up in a meaningful way, then the `list` command can give useful information on how the power shelves in racks are organized:

Example

```
[basecm11->rack]% list
          x-Coord-  y-Coord
Name (key) Room inate   inate Height Devices
-----
A01      T     1       1      48
A02      T     1       2      48
A03      T     1       3      48
A04      T     1       4      48
A05      T     1       5      48      a05-p1-pwr-01..a05-p1-pwr-08,a05...
A06      T     1       6      48
A07      T     1       7      48
A08      T     1       8      48
A09      T     1       9      48
A10     T     1      10      48
A11     T     1      11      48
A12     T     1      12      48
...
...
```

The `help set` command describes how these values can be set.

In the preceding example, the administrator has noted that a rack, for example A05, is in a room T, and is at an (x,y) coordinate of (1,5) in that room.

The power shelf devices have also been given meaningful names: a05-p1-pwr-01, a05-p1-pwr-02..., which here has a prefix that is the rack it is in, and uses `pwr` to indicate it is a power shelf.

4.8.2 The rackoverview Command

The `rackoverview` command displays information about the types of entities in a rack. It also then lists some more detailed information about some of the entity types.

The power shelf one of the types that has more detailed information listed (some output truncated and ellipsized, because this section is focussing on the power shelves):

Example

```
[basecm11->rack]% rackoverview <TAB><TAB>
a01 a02 a03 a04 a05 a06 a07 a08 a09 a10 a11 a12 b01 b02 b03 b04...
[basecm11->rack]% rackoverview a05
Type           Up        Down      Closed      Total
-----
Nodes          18        0         0          18
DPU nodes      0         0         0          0
Managed switches 0        0         0          0
NVLink switches 0        9         0          9
Power shelves   8        0         0          8
Devices         0        0         0          0
Cores          2,592      -        -          2,592
GPUs           72        -        -          72
...
Power          Input    Output      Temperature    Fan      Active    Total
shelf          power   power      Temperature   speed    PSU      PSU
-----
a05-p1-pwr-01  0 W    0 W       0 C        0 RPM    0       6
a05-p1-pwr-02  0 W    0 W       0 C        0 RPM    0       6
a05-p1-pwr-03  3.7 KW 3.6 KW   43.3333 C  6.7 KRPM 6       6
a05-p1-pwr-04  2.51 KW 2.38 KW  42.75 C   6.7 KRPM 6       6
```

```
a05-p1-pwr-05  2.83 KW  2.69 KW  49 C      6.8 KRPM  6      6
a05-p1-pwr-06  3.7 KW   3.5 KW   45 C      6.7 KRPM  6      6
a05-p1-pwr-07  3.7 KW   3.5 KW   46.3333 C  6.7 KRPM  6      6
a05-p1-pwr-08  3.8 KW   3.7 KW   47.3333 C  6.7 KRPM  6      6
[basecm11->rack]%
```

In the preceding, the A05 rack is seen to have some active power shelves, which show some power-related values.

4.8.3 The display Command

If the cluster administrator has recorded the device positions in the rack, then the `display` command in `rack` mode is useful for seeing the position of where power shelves are located in the rack:

Example

```
[basecm11->rack]% display |less -R
...
          A05           B05

48          48
47          47
46          46
45          45
44          44
43          43
42  a05-p1-pwr-08  42  B05-P1-PWR-08
41  a05-p1-pwr-07  41  B05-P1-PWR-07
40  a05-p1-pwr-06  40  B05-P1-PWR-06
39  a05-p1-pwr-05  39  B05-P1-PWR-05
38          38
37  a05-p1-dgx-01-c18 37  b05-p1-dgx-05-c18
36  a05-p1-dgx-01-c17 36  b05-p1-dgx-05-c17
35  a05-p1-dgx-01-c16 35  b05-p1-dgx-05-c16
34  a05-p1-dgx-01-c15 34  b05-p1-dgx-05-c15
33  a05-p1-dgx-01-c14 33  b05-p1-dgx-05-c14
32  a05-p1-dgx-01-c13 32  b05-p1-dgx-05-c13
31  a05-p1-dgx-01-c12 31  b05-p1-dgx-05-c12
30  a05-p1-dgx-01-c11 30  b05-p1-dgx-05-c11
29  a05-p1-dgx-01-c10 29  b05-p1-dgx-05-c10
28  a05-p1-dgx-01-c09 28  b05-p1-dgx-05-c09
27  a05-p1-nvsw-09   27  B05-P1-NVSW-09
26  a05-p1-nvsw-08   26  B05-P1-NVSW-08
25  a05-p1-nvsw-07   25  B05-P1-NVSW-07
24  a05-p1-nvsw-06   24  B05-P1-NVSW-06
23  a05-p1-nvsw-05   23  B05-P1-NVSW-05
22  a05-p1-nvsw-04   22  B05-P1-NVSW-04
21  a05-p1-nvsw-03   21  B05-P1-NVSW-03
20  a05-p1-nvsw-02   20  B05-P1-NVSW-02
19  a05-p1-nvsw-01   19  B05-P1-NVSW-01
18  a05-p1-dgx-01-c08 18  b05-p1-dgx-05-c08
17  a05-p1-dgx-01-c07 17  b05-p1-dgx-05-c07
16  a05-p1-dgx-01-c06 16  b05-p1-dgx-05-c06
15  a05-p1-dgx-01-c05 15  b05-p1-dgx-05-c05
14  a05-p1-dgx-01-c04 14  b05-p1-dgx-05-c04
13  a05-p1-dgx-01-c03 13  b05-p1-dgx-05-c03
12  a05-p1-dgx-01-c02 12  b05-p1-dgx-05-c02
11  a05-p1-dgx-01-c01 11  b05-p1-dgx-05-c01
10          10
```

09	a05-p1-pwr-04	09	B05-P1-PWR-04
08	a05-p1-pwr-03	08	B05-P1-PWR-03
07	a05-p1-pwr-02	07	B05-P1-PWR-02
06	a05-p1-pwr-01	06	B05-P1-PWR-01
05		05	
04		04	
03		03	
02		02	
01		01	

4.8.4 Power Management Of Racks

For devices in general, power operation commands are described in Section 4.2 of the *Administrator Manual*. Power status responses are described on page 220 of the *Administrator Manual*.

Racks can take a while to power up through the configured sequence of components (power shelves, switches, nodes). The BUSY status can therefore be visible for a while. In this status, the power up operation can be cancelled.

A session showing power commands running on the racks of a DGX OS cluster may show output similar to the following:

Example

```
[basecm11->rack]% events off
Private events: off
Broadcast events: off
[basecm11->rack]% power on -b nvl[001-016]
[basecm11->rack]% power status
post::powershelf ..... [ BUSY ] nvl001
post::powershelf ..... [ BUSY ] nvl002
post::powershelf ..... [ BUSY ] nvl003
...
[basecm11->rack]% power status
post::switch ..... [ BUSY ] nvl001
post::switch ..... [ BUSY ] nvl002
post::switch ..... [ BUSY ] nvl003
...
[basecm11->rack]% power cancel nvl001
Cancelled: 1 operations
[basecm11->rack]% power status
none ..... [ CANCELLED ] nvl001
post::compute ..... [ BUSY ] nvl002
post::compute ..... [ BUSY ] nvl003
...
```

More details on power commands use in the `rack` mode of cmsh can be seen with:

Example

```
[basecm11->rack]% help power
Name:
    power - Manipulate the power state of rack

Usage:
    power [OPTIONS] on
    power [OPTIONS] off
    power [OPTIONS] cancel
```

```
power [OPTIONS] status

Options:
  -b, --background
    Run in background, output will come as events

  -d, --delay <seconds>
    Wait <seconds> between executing two sequential power commands. This option is
    ignored for the status command

  -m, --mode <mode>
    Mode: full(default), domain(switch and compute tray), compute, emergency(off only)

  -w, --overview
    Create overview instead of individual rack list

  -f, --force
    Force power command on devices which have been closed or are reported as leaking

  --timeout <seconds>
    Maximal time to wait for all result to come in

  --message <message>
    An optional message saved to DB for on/off/reset operations

  --initial-wait-switch <seconds>
    Wait the specified time before switch power on after power shelves were found UP

  --initial-wait-node <seconds>
    Wait the specified time before node power on after switches were found UP

  --maximal-redfish-wait <seconds>
    Maximal time to wait between power shelf power on and the redfish of switch/node becoming
    available

  --maximum-wait <seconds=30>
    Maximal time to wait between for devices to enter the correct state
```

Examples:

power on rack01	Power on rack01
power off -m compute rack01	Power off the compute tray in rack01
power cancel	Cancel all active rack power operations
power cancel rack01	Cancel power operation on rack01
power status	Show status of all active power operations

5 NVIDIA Autonomous Job Recovery

5.1 Introduction

NVIDIA Mission Control can provide autonomous job recovery. This capability can be installed on a cluster managed by BCM.

Autonomous job recovery can

- monitor Kubernetes jobs and pods
- monitor Slurm jobs and scheduling

and analyze what is being monitored. If it detects processes going wrong, then it can take action to recover from what is wrong, with the aim of reducing downtime.

Autonomous job recovery can display its monitoring visually in a Grafana dashboard using Loki logging.

Autonomous job recovery can be installed and integrated into BCM if Kubernetes and Slurm are installed first.

The physical hardware requirements for a cluster running autonomous job recovery are:

- the nodes should have at least 16 GB of RAM
- the cluster should have at least 30 CPU cores
- the cluster should have at least 64 GB RAM on the head node
- the cluster should have at least 20GB of persistent storage

5.2 Prerequisites To Install Autonomous Job Recovery: Kubernetes Installation

To install autonomous job recovery, Kubernetes should already have been installed and integrated with BCM with appropriate components. Kubernetes installation and integration is typically done with the `cm-kubernetes-setup` utility (section 4.2 of the *Containerization Manual*).

5.2.1 Kubernetes Components Needed For Installation Of Autonomous Job Recovery

The Kubernetes deployment procedure should be carried out so that it includes at least the following components (figure 5.1):

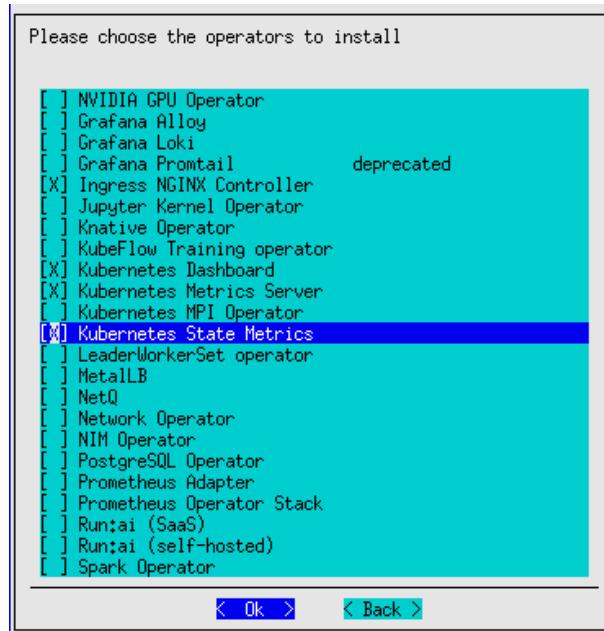


Figure 5.1: Selection of Kubernetes components for installing autonomous job recovery

- Prometheus Operator Stack
- Prometheus Adapter
- Grafana Promtail
- Kubernetes Metrics Server
- Kubernetes State Metrics
- Ingress NGINX Controller

Kyverno should not be installed at the time of writing (April 2025).

5.2.2 Settings For Kubernetes Components Used In Autonomous Job Recovery Installation

In other screens that come up with the preceding components selection, the following settings should be used:

- In the screen asking to configure the Kubernetes StorageClass, a local path storage class should be enabled, with the default NFS storage enabled (figure 5.2):

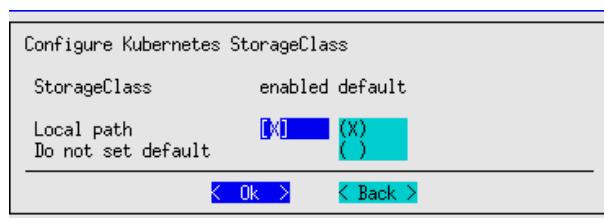


Figure 5.2: Setting the storage class

- In the screen that asks about Loki access (figure 5.3):
 - Loki should be added as the data source for Grafana

- the Loki API should be exposed for Ingress

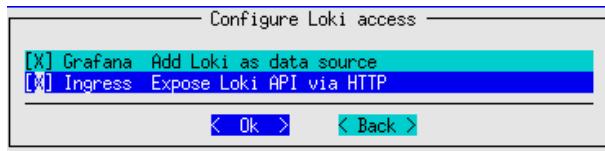


Figure 5.3: Setting up Loki as data source for Grafana, and Loki access via Ingress

- In the screen that configures logging for Grafana (figure 5.4), Grafana Promtail is recommended to have its collection disabled for /var/log from the Kubernetes worker nodes.

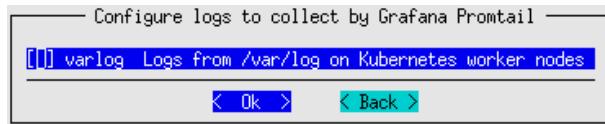


Figure 5.4: Log collection configuration for Grafana Promtail

5.3 Prerequisites To Install Autonomous Job Recovery: Slurm Installation

To install autonomous job recovery, Slurm should already have been installed and integrated with BCM. Slurm installation and integration is typically done with the `cm-wlm-setup` utility (section 7.3.2 of the *Administrator Manual*).

If Slurm is not to be run during regular use, then Slurm should still be installed before autonomous job recovery is installed, despite autonomous job recovery not being applied to Slurm during regular use. This is because autonomous job recovery installation makes use of Slurm configuration.

5.4 Autonomous Job Recovery Installation With `cm-mission-control-setup`

If Kubernetes has been installed (section 5.2), and Slurm has been installed (section 5.3), then autonomous job recovery can be installed with the `cm-mission-control-setup` utility. The `cm-mission-control-setup` utility is part of the `cm-setup` package.

5.4.1 Credentials Required To Install Autonomous Job Recovery When Running `cm-mission-control-setup`

The following credentials are required to install autonomous job recovery:

- Ingress credentials for Loki API access. This would have been set up during the `cm-kubernetes-setup` session.
- MySQL root password for the main CMDaemon database. By default this is the same as the root user password of the operating system.
- NVCR personal token for operator and container images access. The personal token is a string with the prefix: `nvapi-`

5.4.2 Running `cm-mission-control-setup`

The `cm-mission-control-setup` script should be run without options on the active head node. It brings up a TUI dialog (figure 5.5):

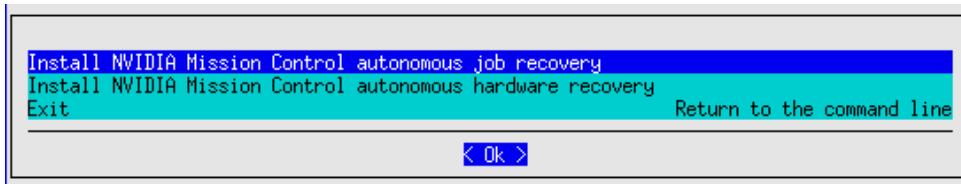


Figure 5.5: Selection from cm-mission-control-setup of autonomous job recovery

If there is more than one Kubernetes instance on the cluster, then a screen appears asking for one of them to be selected for the installation of autonomous job recovery.

If there is more than one Slurm WLM instance on the cluster, then a screen appears asking for one of them to be selected for the installation of autonomous job recovery.

Next, a screen appears asking for the MySQL access credentials.

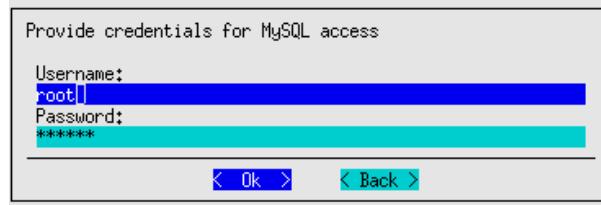


Figure 5.6: Autonomous job recovery deployment configuration: MySQL access

Next, a screen appears asking for the Helm repository credentials.

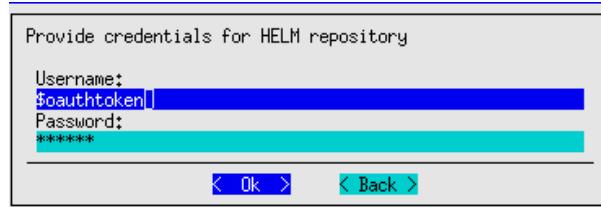


Figure 5.7: Autonomous job recovery deployment configuration: Helm repository access

Next, a screen appears asking for the credentials for Loki API access via Ingress.

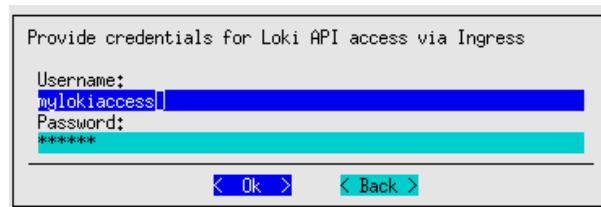


Figure 5.8: Autonomous job recovery deployment configuration: Loki API access via Ingress

The configuration settings can then be saved and deployment of autonomous job recovery is started.

Deployment can take up to 30 minutes.

5.5 Post-installation Checks

5.5.1 Grafana Access

Grafana for the Kubernetes Prometheus stack is exposed by default at:

`https://<cluster IP address>/grafana/login`

as set by the Kubernetes NGINX Ingress Controller.

The default administrator username and password are not changed automatically.

5.5.2 Autonomous Job Recovery Failure When The Slurm Daemon Spool Directory Is Not Cleaned Up

The `slurmd` spool directory is normally cleaned up automatically when a compute node reboots. The `slurmd` spool directory path is defined in the Slurm configuration file `slurm.conf` by the parameter `SlurmdSpoolDir`. `SlurmdSpoolDir` has a default value of `/cm/local/apps/slurm/var/spool/`.

A clean `slurmd` spool is needed after rebooting for Slurm to function correctly. If the spool directory contents persist after a reboot, then it means that a job may not resume properly.

The spool directory contents can persist in, for example, the following use case:

- when autonomous job recovery decides to reboot a node where a job is running, and
- if the job is submitted with the `--requeue` command line option, and
- it uses PMIX

In that use case, Slurm does not clean up PMIX temporary files such as `stepd.slurm.pmix.$SLURM_JOBID.0` from the default spool directory. So when the job starts on the rebooted node, `slurmstepd` cannot start the job again because the file exists.

6 NVIDIA Autonomous Hardware Recovery

6.1 Introduction

NVIDIA Mission Control can provide autonomous hardware recovery. This capability can be installed on a cluster managed by BCM.

It is designed to allow a user to carry out:

- automated baseline tests to validate hardware
- automated health checks to detect hardware failures at the tray, rack, node, and system levels
- automated break/fix workflows (guided recovery on failure)

Autonomous hardware recovery can be integrated with an NVIDIA Base Command Manager (BCM) that is running a Kubernetes instance. The integration is carried out with the BCM script `cm-mission-control-setup` (section 6.3).

The script deploys the autonomous hardware recovery *agents* on all nodes. These agents communicate with the autonomous hardware recovery *backends* that are deployed within Kubernetes worker nodes.

BCM integration allows API access to the agents and provides a GUI to carry out *runbooks*.

More details on autonomous hardware recovery can be found in its dedicated documentation.

The scope of this chapter of the manual is to give guidance on autonomous hardware recovery deployment, and how to access its API and GUI.

6.2 Prerequisites To Install Autonomous Hardware Recovery

Before running the `cm-mission-control-setup` script to install autonomous hardware recovery, the following prerequisites must be met:

- A Kubernetes instance that is integrated with BCM (section 4.2 of the *Containerization Manual*) must be running.
- The Kubernetes instance must have the Kubernetes NGINX Ingress controller add-on for route-based access control installed.
- A user must be added if an extra administrator account is wanted for autonomous hardware recovery.
For example:

Example

```
root@basecm11:~# cmsh
[basecm11]% user add extraadmin
[basecm11->user*[extraadmin*]]% set password
enter new password:
retype new password:
[basecm11->user*[extraadmin*]]% commit
```

This extra administrator can then be specified when installing with `cm-mission-control-setup` later on (section 6.3).

- At the time of writing of this section (August 2025), Kyverno must not be installed on the Kubernetes instance, or the installation fails. This is expected to change later on.

As a rough guide to the physical hardware requirements for autonomous hardware recovery, it is expected that a cluster with hundreds of nodes should use a backend node with at least 16 CPU cores per node, at least 32 GB RAM per node, and at least 2 local (non-NFS) hard drives. The hard drives should be at least:

- one unpartitioned drive of 1.5 TB
- one drive with filesystems
 - one filesystem with at least 500 GB available for the autonomous hardware recovery backend
 - one filesystem with at least 20 GB available under `/var` for storing container images

More precise physical hardware requirements are best determined with the help of NVIDIA engineers.

6.3 Installation Of NVIDIA Autonomous Hardware Recovery Using `cm-mission-control-setup`

Running `cm-mission-control-setup` as the root user starts up the TUI installation (figure 6.1):

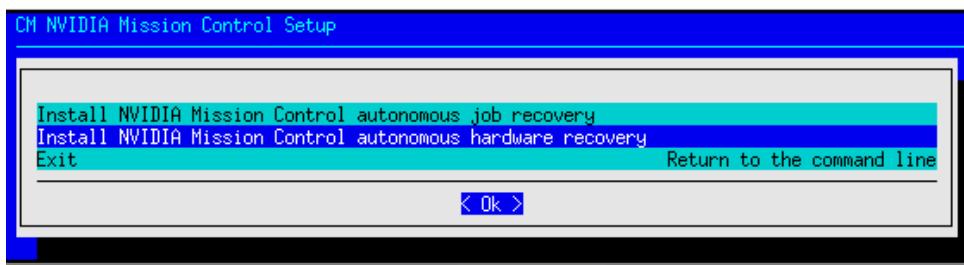


Figure 6.1: Selection from `cm-mission-control-setup` of autonomous hardware recovery

Selecting `NVIDIA Mission Control autonomous hardware recovery` brings up a user selection screen. The extra administrator account(s) set up earlier for administration of autonomous hardware recovery can then be specified.

After the extra administrator accounts have been specified, the Helm registry credentials screen (figure 6.2) comes up:

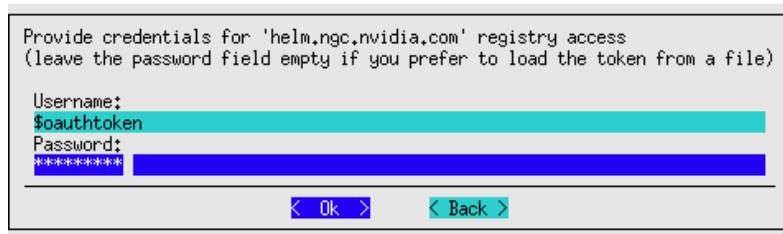


Figure 6.2: Setting Helm credentials for autonomous hardware recovery

If the Helm registry password is accepted, then the NVCR registry credentials screen (figure 6.3) comes up:

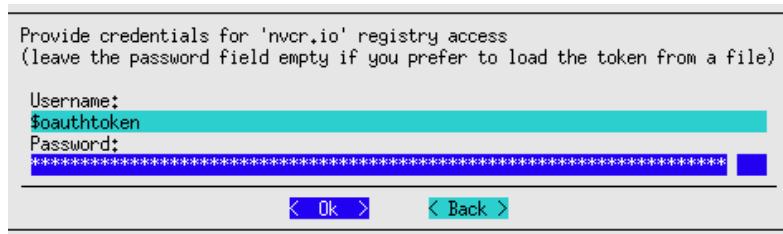


Figure 6.3: Setting NVCR credentials for autonomous hardware recovery

If the NVCR registry password is accepted, then the API registry credentials screen (figure 6.4) comes up:

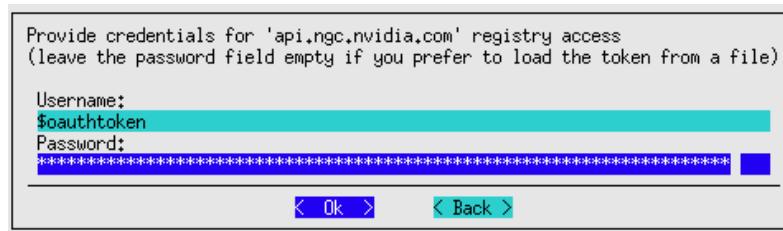


Figure 6.4: Setting API credentials for autonomous hardware recovery

If the password for API registry access is accepted, then a TLS certificates screen (figure 6.5) comes up:

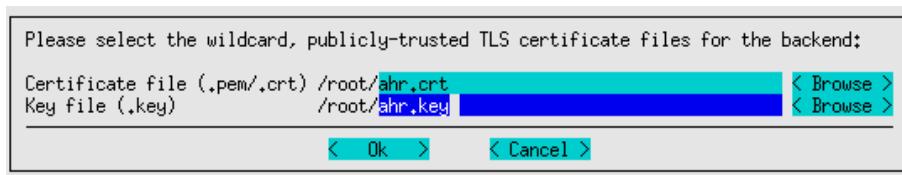


Figure 6.5: Setting TLS certificate files for autonomous hardware recovery

After the path to the TLS certificates are entered, a backend node selection screen (figure 6.6) comes up:

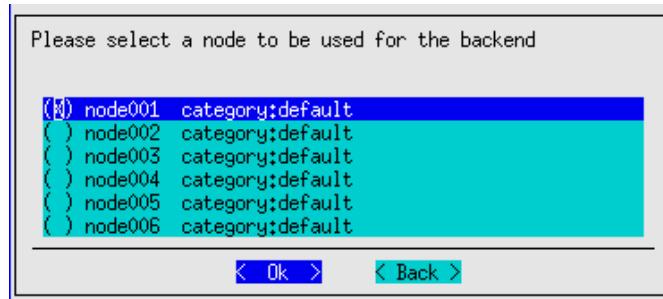


Figure 6.6: Setting a backend node for autonomous hardware recovery

After a backend node is selected, a failover node selection screen (figure 6.7) comes up. Selection is optional:

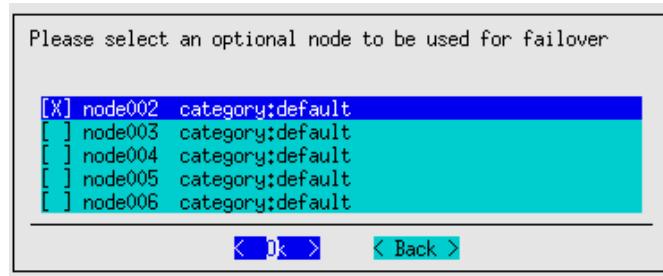


Figure 6.7: Setting an optional failover node for autonomous hardware recovery

After that an agent category selection screen (figure 6.8) comes up:

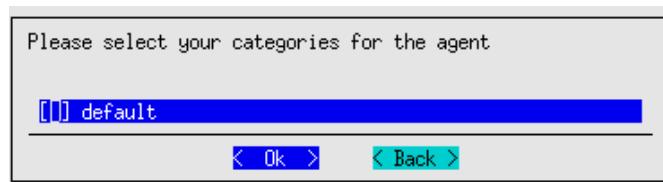


Figure 6.8: Setting an agent category for autonomous hardware recovery

After that a backend customizations screen (figure 6.9) comes up:

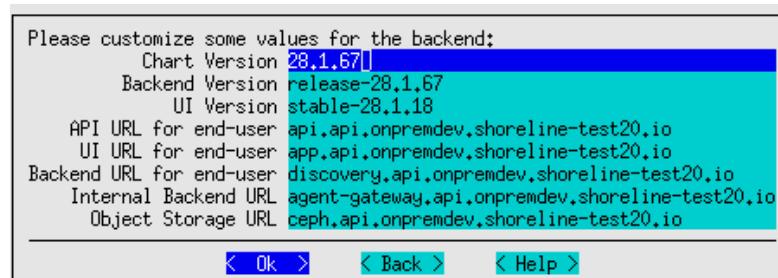


Figure 6.9: Setting backend customizations for autonomous hardware recovery

The endpoints set in the backend should resolve either to the active head node, or to the backend pod node.

A warnings screen appears if the domains cannot be resolved (figure 6.10):

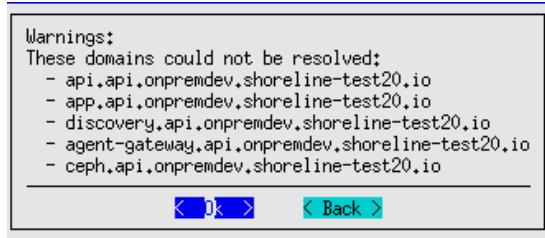


Figure 6.10: Configuration warnings for autonomous hardware recovery

After that a storage customizations screen for the backend appears (figure 6.11) with suggested default values:

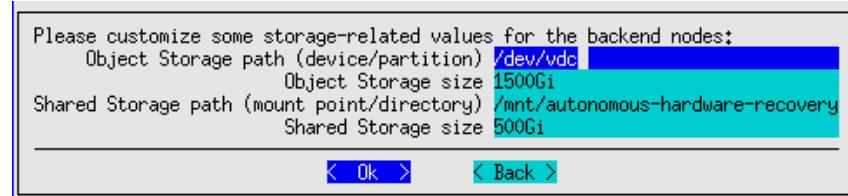


Figure 6.11: Storage customizations for the backend for autonomous hardware recovery

The TUI then asks if monitoring should be set up.

If monitoring is to be configured, then values for the Prometheus and Grafana Tempo endpoints, and for the Prometheus and Grafana Tempo API keys must be entered. The HTTP user ID and password for Loki logging must also be entered (figure 6.12):

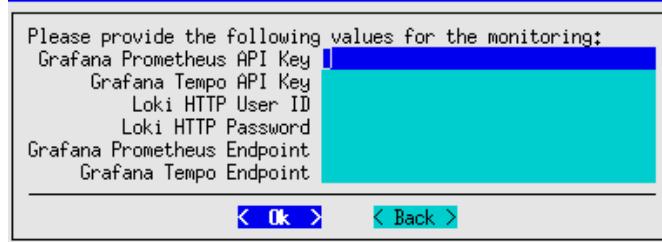


Figure 6.12: Monitoring values configuration for autonomous hardware recovery

The agent version should be set (figure 6.13):

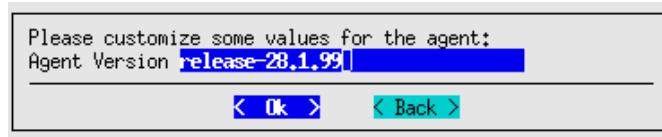


Figure 6.13: Agent version setting for autonomous hardware recovery

The configuration from the wizard can be saved and the configuration can be deployed (figure 6.14):

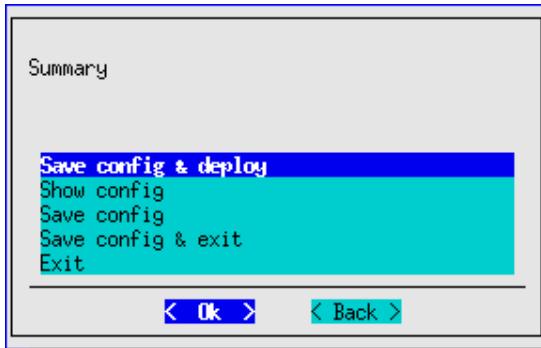


Figure 6.14: Saving the configuration and deploying it for autonomous hardware recovery

6.4 Verifying Autonomous Hardware Recovery Is Up And Ready

6.4.1 Checking The Pods

The pods that run autonomous hardware recovery should be visible in their namespace after the deployment. Their status should be Running, or Completed:

Example

```
root@basecm11:~# kubectl get pods -n autonomous-hardware-recovery
NAME                      READY   STATUS    RESTARTS   AGE
shoreline-backup-29208910-dwrvp  0/1     Completed  0          14m
shoreline-backup-29208915-6fghc  0/1     Completed  0          9m19s
shoreline-backup-29208920-17h25  0/1     Completed  0          4m19s
shoreline-disk-checker-8tsl4    0/1     Completed  0          47m
shorelinebackend-0             7/7     Running   0          46m
shorelinebackend-failover-0    7/7     Running   0          46m
root@basecm11:~#
```

6.4.2 Testing The Web Interface

The API is described in <https://ssapi.shorelinesoftware.net/>.

Access to the UI website can be tested by running a curl command against the home endpoint page of its domain. The domain was earlier picked up from its certificates, and in the following example is app.api.onpremdev.shoreline-test20.io

Example

```
root@basecm11:~# curl https://app.api.onpremdev.shoreline-test20.io/home/
```

This should get output similar to:

Example

```
<!doctype html>
<html lang="en">
  <head>
    <base href="/home/" />

    <meta charset="UTF-8" />
    <meta name="viewport" content="width=device-width, initial-scale=1.0" />

    <script>
      window.$$version = 'stable-28.1.25';
```

```
</script>
<script type="module" crossorigin src="/home/assets/index-CMnWoHWq.js"></script>
<link rel="stylesheet" crossorigin href="/home/assets/index-DEW7GL4G.css">
</head>
<body>
  <div id="app"></div>
</body>
</html>
```

The web UI can be accessed if the API domain and the application domain point to the external IP address of the head node.

If for some reason DNS resolution is not yet set up for the domain, then, for example, if the head node is at the IP address 10.3.195.201, the following entries can be added to /etc/hosts on the head node for the api and app subdomains:

```
10.3.195.201 api.api.onpremdev.shoreline-test20.io
10.3.195.201 app.api.onpremdev.shoreline-test20.io
```

A web browser that accesses <https://app.api.onpremdev.shoreline-test20.io/> from the head node then displays the autonomous hardware recovery interface landing page (figure 6.15):

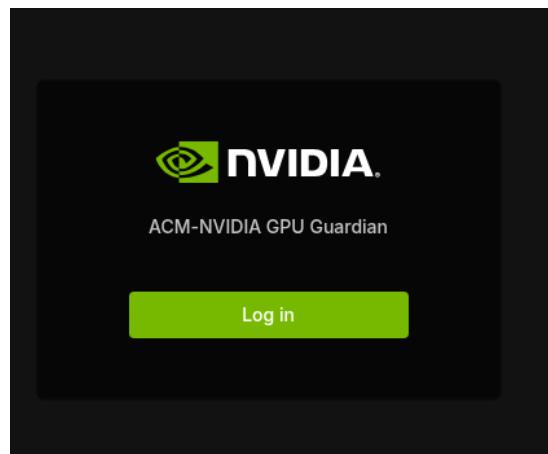


Figure 6.15: Web GUI for autonomous hardware recovery: landing page

A login can be carried out on the GUI with a username/password authentication (figure 6.16):

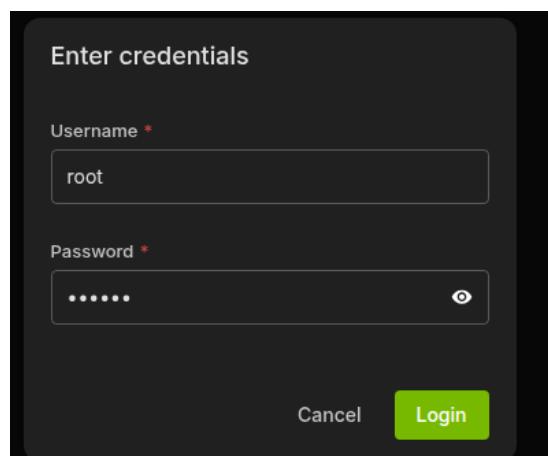


Figure 6.16: Web GUI for autonomous hardware recovery: login page

This opens up the autonomous hardware recovery dashboard.

A `hello world` runbook can be executed after login as follows:

- The runbook window is opened using the navigation panel item Runbooks
- A new runbook is created by clicking on the New Runbook button, and then on the op statement option (figure 6.17):

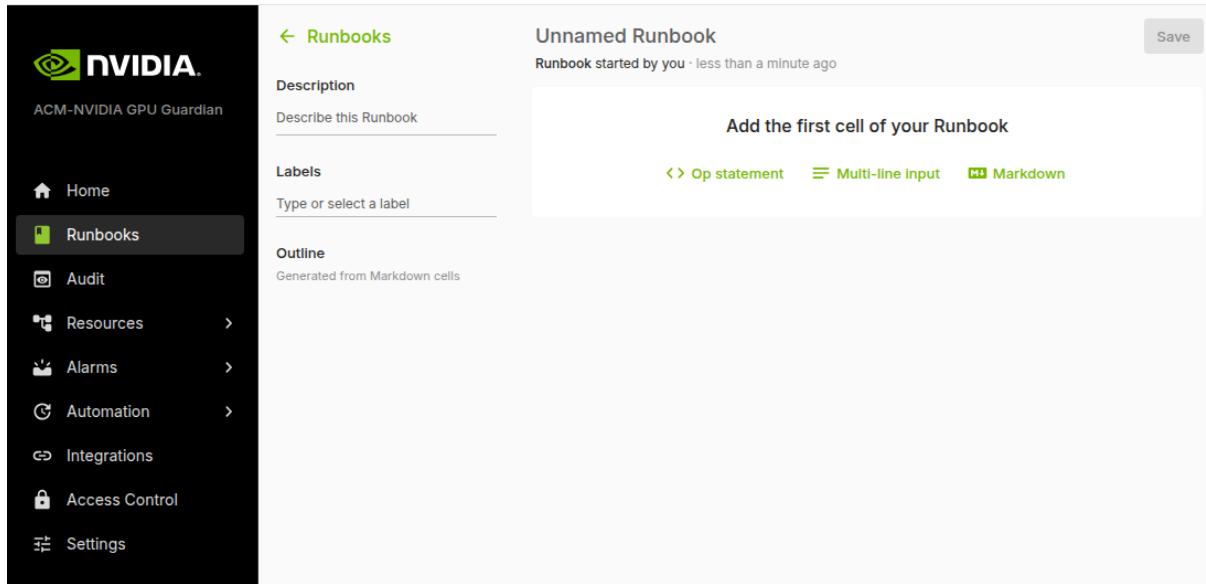


Figure 6.17: Web GUI for autonomous hardware recovery: creating a new runbook

- A cell opens up and prompts the user to enter a statement.

The following statement can be added:

```
host | `echo Hello, World!`
```

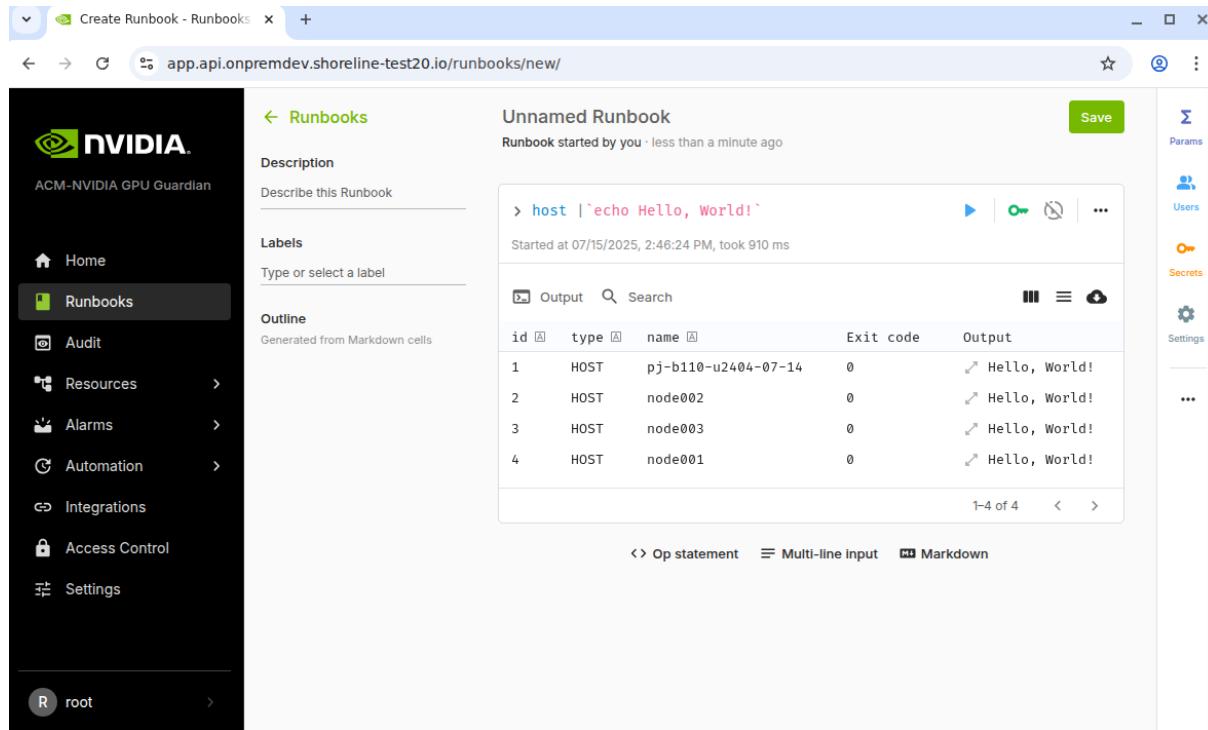


Figure 6.18: Web GUI for autonomous hardware recovery: adding a “hello world” runbook

- The runbook can be executed with the enter key or by clicking on the Execute button next to the cell. There is hovertext to explain the meanings of the widgets.
Outputs of Hello, World! should show up in the GUI (figure 6.18):
- The runbook can be saved for posterity with the Save button.

If the preceding runbook session works, then it indicates that autonomous hardware recovery has been deployed properly.

7 NVLink Switch Integration

The NVIDIA GB200 server platforms have 9 NVLink switch trays per rack. An NVL72 reference rack configuration is described in, for example, [the NVIDIA DGX SuperPOD configuration](#).

The switch trays connect GPUs in each rack (36 per rack, or 72 per rack) to each other on a 400 Gb/s InfiniBand NDR network.

This chapter considers an NVIDIA GB200 server platform that has been configured to work with BCM.

7.1 NVLink Switch Listing And Overview

The `list -t switch` command in device mode can be used to list all switches of a configured cluster. The `list -field kind=nvlink` command in device mode can be used to list only the NVLink switches of a configured cluster.

NVLink switches are configured as devices on an BMC network, as indicated by the following example (some rows ellipsized):

Example

```
[maple->device]# list -t switch
Type      Hostname (key)      MAC          IP           Network      Status
-----  -----
Switch    dgx-gb200-m06-nvlsw1  E0:9D:73:05:70:6C  10.140.0.45  ipminet0  [  UP   ]
Switch    dgx-gb200-m07-nvlsw1  B8:E9:24:B9:84:8C  10.140.0.48  ipminet0  [  UP   ]
Switch    dgx-gb200-n01-nvlsw1  E0:9D:73:05:70:5C  10.140.0.51  ipminet0  [  UP   ]
Switch    dgx-gb200-n07-nvlsw1  E0:9D:73:05:70:AC  10.140.0.54  ipminet0  [  UP   ]
...
...
```

Example

```
[basecm11->device[dgx-gb200-m06-nvlsw1]]% switchoverview
Device      : dgx-gb200-m06-nvlsw1
State       : [  UP   ]
Model       : N5400_LD
Serial number : MT1234556789
Part number  : 123-4567-89

Port assignment:

Port  Name  Status Uplink Speed      RX      TX
-----  -----
1     acp1  UP     no    400 Gb/s  64 TiB  64 TiB
```

2	acp2	UP	no	400 Gb/s	64 TiB	64 TiB
...						
143	acp143	UP	no	400 Gb/s	64 TiB	64 TiB
144	acp144	UP	no	400 Gb/s	64 TiB	64 TiB
145	eth0	UP	no	1.00 Gb/s	1.22 TiB	6.0 GiB
146	eth1	DOWN	no	0 b/s	4.5 KiB	28.7 KiB
147	fnm1	UP	no	100 Gb/s	0 B	0 B
148	fnm2	UP	no	100 Gb/s	0 B	0 B

7.2 NVLink Switch Configuration

Each NVLink switch has two redundant ethernet interfaces and a Redfish interface.

Example

```
[basecm11->device[a05-p1-nvsw-01]->interfaces]% list
Type           Network device name   IP                  Network
-----
bmc            rf0                7.241.3.21        ipminet2
physical       eth0               7.241.3.1        ipminet2
physical       eth1               7.241.3.11       ipminet2
```

An NVLink runs NVOS and can be configured to boot via ZTP. The JSON template ZTP file that is provided by BCM is used to check the configured image, and to install the cm-lite-daemon. If ZTP is configured correctly before the switch first boots, then it is ready for use.

Example

```
[basecm11->device[a05-p1-nvsw-01]->ztpsettings]% show
Parameter          Value
-----
Revision
Script template    nvos-ztp.sh
JSON template      nvlink-nvos.json
Image              nvos-amd64-25.02.2141.bin
Check image on boot no
Run ZTP on each boot yes
Install lite daemon yes
Enable API         yes
```

On an NVLink switch that is already booted, or where ZTP is disabled, cm-lite-daemon can be installed manually.

The first step is to download the relevant packages using cmsh.

The download needs to be executed once for all switches. If updates are available in the BCM repositories, then the download should be repeated and the subsequent installation process should also be repeated.

```
[basecm11%] litedaemon download
Download all/cm-config-cm_1trunk_all.deb for 2004: done
Download all/cm-lite-daemon_1master-100775-cm-752a72d739_all.deb for 2004: done
Download amd64/cm-openssl_3.1.8-100204-cm-a2b95b5141_amd64.deb for 2004: done
Download amd64/cm-python312_3.12.11-100035-cm-f4ae26a924_amd64.deb for 2004: done
Download amd64/cm-python3_1master-100155-cm-eaf713f03f_amd64.deb for 2004: done
...
```

After the packages are downloaded, they can be installed on the NVLink switches using the litedaemon package manager with the `install` command. It is recommended to do all 9 switches for a rack at the same time; multiple racks can be done at once.

```
[basecm11->device]% litedaemon install --rack a05 --intersection --field kind=nvlink
success: yes
--- stdout ---
run: /home/admin/cm-lite-daemon-install.py --head-ips 10.141.255.254 --port 8081
    --hostname a05-p1-nvsw-01 --interface eth0
+ ping -c 1 10.141.255.254
Using IP 10.141.255.254 without supplied vrf
* Download: https://10.141.255.254:8081/switch/packages/2404/files *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-config-cm_1trunk_all.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-lite-daemon_1master-
    100775-cm-752a72d739_all.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-openssl_3.1.8-100204-
    cm-a2b95b5141_amd64.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-python312_3.12.11-
    100035-cm-f4ae26a924_amd64.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-python3_1master-100155-
    cm-eaf713f03f_amd64.deb *
* Install *
* Setup *
* Setup from /home/admin *
* Certificates *
- cluster.pem
- bootstrap.pem
- bootstrap.key
* Register *
* Done *
```

When updated packages have been downloaded they also need to be updated on the NVLink switches. Updates can be carried out using the --update option with `litedaemon install` as follows:

```
[basecm11->device]% litedaemon install --update --field kind=nvlink
run: /home/admin/cm-lite-daemon-install.py --head-ips 10.141.255.254 --port 8081
    --hostname a05-p1-nvsw-01 --interface eth0 --update
+ ping -c 1 10.141.255.254
Using IP 10.141.255.254 without supplied vrf
* Download: https://10.141.255.254:8081/switch/packages/2404/files *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-config-cm_1trunk
    _all.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-lite-daemon_1master-
    100775-cm-752a72d739_all.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-openssl_3.1.8-100204-
    cm-a2b95b5141_amd64.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-python312_3.12.11-
    100035-cm-f4ae26a924_amd64.deb *
* Download: https://10.141.255.254:8081/switch/packages/2404/cm-python3_1master-
    100155-cm-eaf713f03f_amd64.deb *
* Install *
* Restart *
```

Removing all BCM-related packages can be carried out using the --update option with `lightdaemon remove` as follows:

```
[basecm11->device]% litedaemon remove --update --field kind=nvlink
* Systemctl stop *
* Remove *
* Systemctl reload *
```

7.3 NVLink Switch Monitoring

NVLink switches are monitored in two ways.

- OS data using cm-lite-daemon running on the switch
- RedFish sampled from the active head node (currently only leak information is reported)

All data can be viewed using standard BCM monitoring

```
[basecm11->device[a05-p1-nvsw-01]]% latestmonitoringdata
...
LoadOne                      3.21   33s
RF_LD_LeakDetection          OK      1s   Rollup: OK
RF_LeakDetector    leakage1    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage2    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage3    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage4    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage5    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage6    OK      1s   Detector: OK, Health: OK, State: Enabled
RF_LeakDetector    leakage_aggr OK      1s   Detector: OK, Health: OK, State: Enabled
```

7.4 NVLink Switch Redfish Events

NVLink switches report leaks to all tools that are subscribed to Redfish events. BCM does this for all NVLink switches if the Redfish IP address and BMC settings are configured correctly.

The `redfishsubscription` command lists the devices that have a Redfish subscription (page 70).

```
[basecm11->device[a05-p1-nvsw-01]]% redfishsubscription
hostname      id      ip           kind      port      subscribed
-----+-----+-----+-----+-----+-----+
a05-p1-nvsw-01 1234  7.241.3.21  GB200SW  443       true
```

Testing of the event subscription work can be done from inside cmsh.

```
[basecm11->device[a05-p1-nvsw-01]]% redfishevent --message Test
...
```

8 Power Reservation Steering

8.1 Introduction

Power reservation steering (PRS) is about automated speculative planning of power across groups of nodes used by Slurm. PRS is typically done for larger clusters, where minimising the number of idling nodes typically results in significant savings.

The planning is done over a period of time spent measuring existing resource consumption and observing the power consumed. The patterns that are seen in that period of time suggest a direction for the power allocation on the cluster (hence the word *steering*) for the next period of time. Based on the planned steering, a power cap is set for some groups of nodes (power domains) in the cluster. Setting the power cap for those power domains, means that a *reservation* of the remaining power is made for the rest of the cluster during the upcoming period of time.

PRS works using a PRS daemon. This tracks the power availability and manages power planning. A component is the PRS controller that connects to CMdaemon, and which during the time loop period controls the power domains.

DPS (Domain Power Service) is a datacenter level software intended to optimally manage power allocation across the entire datacenter. By default, if a Slurm topology has PRS enabled, then DPS configures PRS with a power budget for nodes.

8.2 Deployment Of PRS

After a Slurm cluster has been configured and deployed (section 7.3.2 of the *Administrator Manual*), PRS can be deployed.

Deployment of PRS is started in BCM by running the `cm-prs-setup` script. This brings up a TUI wizard (figure 8.1):



Figure 8.1: Starting deploying with `cm-prs-setup`

Categories can be selected for PRS (figure 8.2):

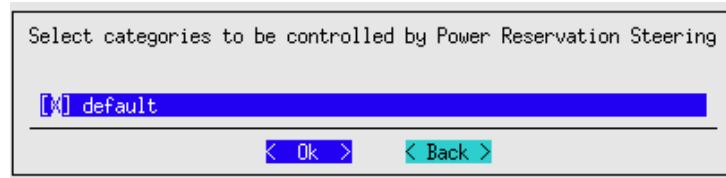


Figure 8.2: Selecting a category with cm-prs-setup

If selecting categories is too inclusive, then individual nodes can be selected instead in a following screen.

After categories or individual nodes have been selected, one or more power domains can be added (figure 8.3):

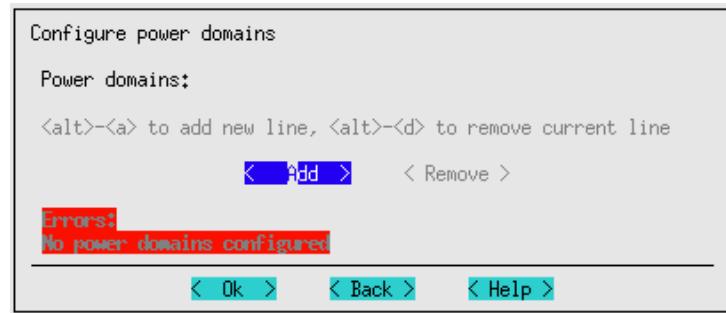


Figure 8.3: Setting a power domain with cm-prs-setup

Each power domain can be given an arbitrary name, and a domain grouping (figure 8.4):

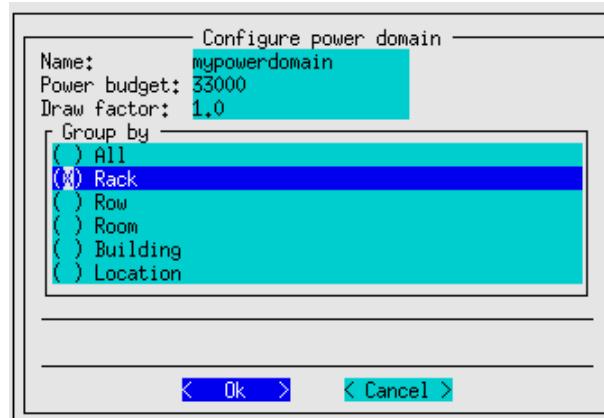


Figure 8.4: Configuring the power domain budget and grouping with cm-prs-setup

The default domain group of **all** is simply one big power domain. Groups such as **rack**, **row**, **room**, and so on become new domains, each with the specified power budget.

The power budgets for the domains are set automatically if the nodes have been powered up. Otherwise the budgets must be specified manually.

The maximum power consumption per node is:

- 400W for A100 nodes
- 700W for H100 nodes
- 1000W for B200 nodes

- 2700W for GB200 nodes

The `Static power` is the power consumption when the node is shut down, and only the BMC is standing by. Typically it is about 20W.

The power draw factor can be set in the range from 0.0-1.0, in steps of 0.1.

The power consumption limits can then be set for the domain.

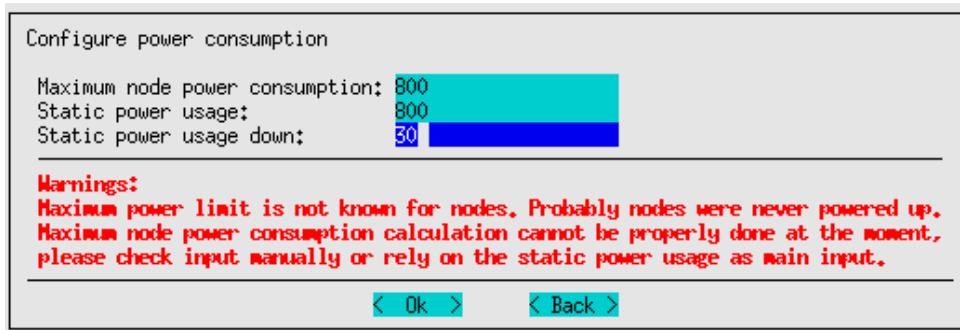


Figure 8.5: Setting power consumption limits with `cm-prs-setup`

After the configuration steps are completed, the configuration can be saved and PRS deployed (figure 8.6):

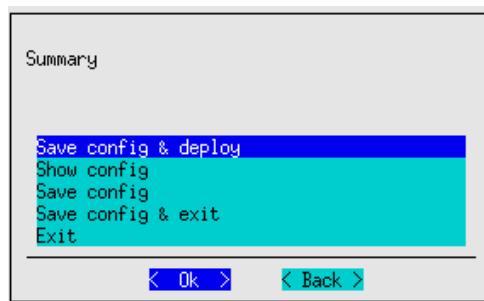


Figure 8.6: Saving the configuration and deploying it with `cm-prs-setup`

8.3 Changes Made On Deployment, And Managing Changes Post-Deployment With cmsh

8.3.1 PRS Values

PRS deployment creates the following configuration overlays:

- `prs-server`: this has the `PRSServerRole` with `PRSDomain` configurations
- `prs-client`: this has the `PRSCClientRole` with power usage configuration

The PRS server configuration overlay is always assigned to the head nodes, and the client overlay is assigned to the nodes and categories configured in the wizard.

PRS server configuration values can be managed in `cmsh` with their role under the configuration overlay:

Example

```
[basecm11->configurationoverlay[prs-server]->roles[PRS::Server]]% show
Parameter          Value
-----
Name              PRS::Server
Revision
Type              PRSServerRole
Add services      yes
Config server port 8880
Job scheduler server port 8881
Timeout           10s
Interval          30s
Window            5
PRS server certificate path /cm/local/apps/prs/etc/server.pem
PRS server private key path /cm/local/apps/prs/etc/server.key
PRS server CA certificate path /cm/local/apps/prs/etc/ca.pem
PRS server CA private key path /cm/local/apps/prs/etc/ca.key
CMD certificate path /cm/local/apps/prs/etc/cmd.pem
CMD private key path /cm/local/apps/prs/etc/cmd.key
PRS client certificate path /cm/local/apps/prs/etc/prs.pem
PRS client private key path /cm/local/apps/prs/etc/prs.key
PRS client CA certificate path /cm/local/apps/cmd/pythoncm/lib/python3.12/site-packages/pythoncm/etc/cacert.pem
Domains           <2 in submode>
[basecm11->configurationoverlay[prs-server]->roles[PRS::Server]]% domains
[basecm11->configurationoverlay[prs-server]->roles[PRS::Server]->domains]% list
Name (key)        Power budget Power budget model Power draw model Power draw factor Group by
-----
anotherdom       30.0KW     scalar        linear        1           all
mypowerdomain   10.0KW     scalar        linear        1           all
[basecm11->configurationoverlay[prs-server]->roles[PRS::Server]->domains]%
```

Similarly, the PRS client role values can be managed under their role:

Example

```
[basecm11->configurationoverlay[prs-client]->roles[PRS::Client]]% show
Parameter          Value
-----
Name              PRS::Client
Revision
Type              PRSClientRole
Add services      yes
Static power usage 800W
Static power usage down 30W
Min CPU power limit 0W
Max CPU power limit 0W
Min GPU power limit 0W
Max GPU power limit 0W
```

8.3.2 Slurm Values

On deployment of PRS, the selectType of the Slurm instance is changed from:

```
select/cons_tres
to
select/gnl_cons_tres
```

Example

```
[basecm11->wlm[slurm]]% get selecttype
select/gnl_cons_tres
```

The certificate paths for PRS for the Slurm instances are managed within the PRS settings of the Slurm instances:

Example

```
[basecm11->wlm[slurm]->prssettings]# show
Parameter          Value
-----
Certificate path    /cm/local/apps/prs/etc/slurm-slurm.pem
Revision           1
Private key path   /cm/local/apps/prs/etc/slurm-slurm.key
```

8.3.3 PRS Removal

The cm-prs-setup script can be used to remove existing PRS domains.

8.4 Changes Made On Deployment, And Managing Changes Post-Deployment With Base View

8.4.1 Base View Power Reservation Steering Wizard

The Base View GUI PRS wizard takes the same inputs as the cm-prs-setup TUI PRS wizard of section 8.3.

The Base View wizard can be accessed using the navigation path:

Mission Control > Power Reservation Steering > Start Wizard

The navigation path is only available if Slurm has previously been installed. Slurm can be installed using the navigation path:

HPC > Workload Management Wizard

The Base View PRS wizard, on starting up, displays an introduction screen as in figure 8.7:

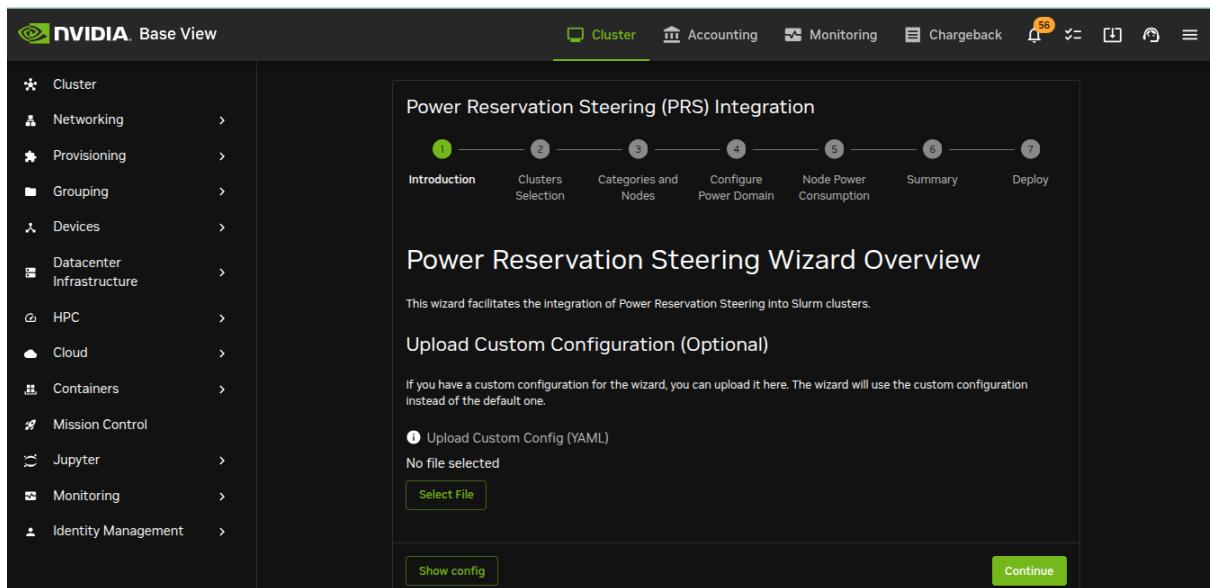


Figure 8.7: Running The PRS Installation Wizard From Base View

The steps that the PRS wizard goes through can be followed in the breadcrumb trail. The steps are:

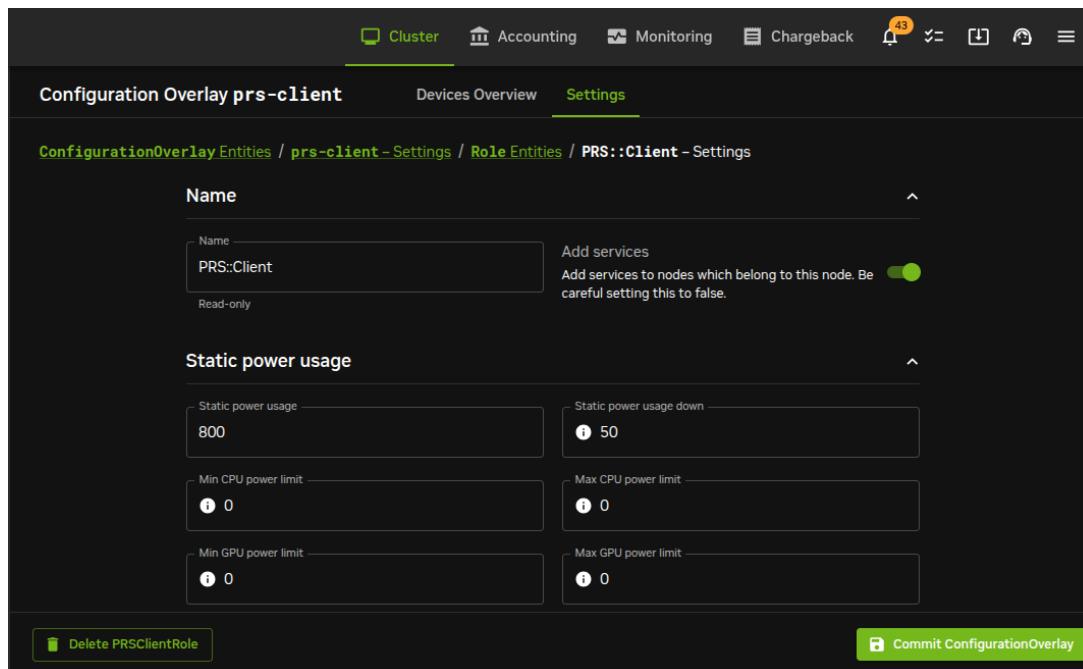
- Introduction: a custom YAML configuration can be uploaded here, and the existing YAML configuration can be viewed
- Clusters Selection: Slurm instances can be selected here for PRS
- Categories and Nodes: regular and head nodes can be chosen for PRS
- Configure Power Domain: a power domain can be set, with settings as in the TUI version (figure 8.4)
- Node Power Consumption: power consumption for the domain can be set
- Summary: a summary of the settings is displayed
- Deploy: the deployment stages are carried out and shown

8.4.2 Base View Power Reservation Steering Client And Server Configuration Overlays

With Base View, the values set for the PRS client and PRS server can be accessed using their roles at configuration overlay, category, and device level.

For example, for the PRS client as configured in the Base View PRS configuration overlay, the navigation path to the configuration pane (figure 8.8) is:

Configuration Overlay > prs-client  > Settings > Roles  > PRS-client 



The screenshot shows the 'Configuration Overlay prs-client' page in the OpenShift web console. The 'Settings' tab is active. The 'Name' field contains 'PRS::Client'. Under 'Static power usage', the 'Static power usage' input is set to 800, and the 'Static power usage down' input is set to 50. Under 'Power limits', both 'Min' and 'Max' values for CPU and GPU power limits are set to 0. At the bottom, there are 'Delete PRSClientRole' and 'Commit ConfigurationOverlay' buttons.

Figure 8.8: PRS Client Values Using The Configuration Role In Base View

9 NVIDIA Mission Control DGX GB200 Measurables

Measurables are metrics, health checks, and enummetrics (appendix G of the *Administrator Manual*).

This section describes the extra DGX GB200 measurables and their parameters that may not have been described in appendix G of the *Administrator Manual*. The DGX GB200 measurables are grouped in this chapter according to the following themes:

- circuit-related metrics (section 9.1)
- leak detection measurables (section 9.2)
- NVLink measurables (section 9.3)
- power shelf measurables (section 9.4)
- cooling distribution unit metrics (section 9.5)
- GPU measurables (section 9.6)
- Prometheus metrics (section 9.7)
- Redfish metrics (section 9.8)
- Redfish enums (section 9.9)
- health checks (section 9.10)

9.1 Circuit-Related DGX GB200 Metrics

Electrical power in the data center is supplied to the DGX GB200 racks from a circuit that is typically from a remote power panel (RPP) or an overhead busway.

Circuit-related metrics for the DGX GB200 platform are shown in table 9.1.

Table 9.1: Circuit-related DGX GB200 metrics

DGX GB200 Circuit Metric	Parameter	Description
CircuitCurrent		
CircuitPhaseCurrent	phase=1	
CircuitPhaseCurrent	phase=2	
CircuitPhaseCurrent	phase=3	
CircuitPower		
RackCircuitCurrent*	circuit=RPP-B12-3	
RackCircuitCurrentLimit*	circuit=RPP-B12-3	
RackCircuitPhaseCurrent*,**	circuit=RPP-B12-3;phase=3	
RackCircuitPower*	circuit=RPP-B12-3	
TotalCircuitCurrent		Total
TotalCircuitCurrentLimit		Total
TotalCircuitPhaseCurrent		Total
TotalCircuitPhaseCurrent	phase=1	Total
TotalCircuitPhaseCurrent	phase=2	Total
TotalCircuitPhaseCurrent	phase=3	Total
TotalCircuitPower		Total

* The parameter set here depends on the listed power circuit name (section 4.7)

** The phase value can be 1, 2, or 3

9.2 Leak Detection DGX GB200 Measurables

Liquid cooling is used in the DGX GB200. The ability to detect leaking coolant is an important enough concern that there are several measurables associated with it.

Leak detection measurables for the DGX GB200 platform are shown in table 9.2:

Table 9.2: Leak detection DGX GB200 measurables

DGX GB200 leak measurable	Parameter	Description
LeakResponseRackElectrical\IsolationStatus		
LeakResponseRackLiquidIso\lationStatus		
LeakSensorFaultRack		
DevicesWithLeaks	Total	Number of devices that have detected a leak
RF_Chassis_0_LeakDetector_0_ColdPlate*		
RF_Chassis_0_LeakDetector_0_Manifold*		
RF_Chassis_0_LeakDetector_0_state		Enum, not metric

...continues

Table 9.2: Leak detection DGX GB200 measurables...continued

DGX GB200 leak measurable	Parameter	Description
ColdPlate*		
RF_Chassis_0_LeakDetector_0_Manifold*	state	Enum, not metric
RF_LD_LeakDetection		Enum
RF_LeakDetector	Chassis_0_LeakDetector_0_ColdPlate*	Enum
RF_LeakDetector	Chassis_0_LeakDetector_0_Manifold*	Enum

* The substring LeakDetector_0 can also take the value LeakDetector_1

9.3 DGX GB200 NVLink Measurables

NVLink-related measurables for the DGX GB200 platform are shown in table 9.3.

The measurables without the RF_ prefix are provided by the ClusterTotal data producer. The measurables with the RF_ prefix are provided via Redfish.

BCM picks up the Redfish measurables and their descriptions automatically. Many measurables still have no associated descriptions. For now, the names of the measurables can be a guide on what they are about.

Table 9.3: DGX GB200 NVLink measurables

DGX GB200 NVLink measurables	Parameter	Description
NVLinkSwitchesClosed		Number of NVLink switches not marked as UP or DOWN
NVLinkSwitchesDown		Number of NVLink switches marked as DOWN
NVLinkSwitchesTotal		Total number of NVLink switches
NVLinkSwitchesUp		Number of NVLink switches marked as UP
RF_NVLink_Port_Nvidia_BitErrorRate*		
RF_NVLink_Port_Nvidia_LinkDowned*		
RF_NVLink_Port_Nvidia_RXNoProtocol*		
RF_NVLink_Port_Nvidia_TXNoProtocol*		
RF_NVLink_Port_Nvidia_TXWait*		
RF_NVLink_Port_RX*		
RF_NVLink_Port_RXFrames_Networking*		
RF_NVLink_Port_RX_Total		
RF_NVLink_Port_TX*		
RF_NVLink_Port_TXFrames_Networking*		

...continues

Table 9.3: DGX GB200 NVLink measurables...continued

DGX GB200 NVLink measurables	Parameter	Description
RF_NVLink_Port_TX_Total		
RF_NVLink_ResourceHealthRollup_Status		
RF_NVLink_ResourceHealth_Status		
RF_NVLink_ResourceState_Status*		
RF_NVLink_Resource_CurrentSpeed*		
RF_NVLink_Resource_LinkState*		
RF_NVLink_Resource_LinkStatus*		
RF_NVLink_Resource_MaxSpeed*		
RF_NVLink_Resource_Nvidia_RXWidth*		
RF_NVLink_Resource_Nvidia_TXWidth*		

* The parameter here is one of 17 ports from the node to the switch.

It can take a value of acp0 to acp17

9.4 DGX GB200 Power Shelf Measurables

The list of power shelf measurables on a DGX GB200 platform are displayed in table 9.4.

These are produced by the `redfish_power_shelf`, `ClusterTotal`, and `AggregatePowerShelf` data producers.

Table 9.4: Power shelf measurables

Power Shelf Measurable	Parameter	Description
AveragePowerShelfTemperature*		Average power shelf temperature
PowerShelvesClosed**		Number of power shelves not marked as UP or DOWN
PowerShelvesDown**		Number of power shelves marked as DOWN
PowerShelvesTotal**		Total number of power shelves
PowerShelvesUp**		Number of power shelves marked as UP
RF_Active_PSU		Number of PSUs of power shelf active
RF_Health_PSU†		Are all PSUs of power shelf active?
RF_PowerShelf_Status [†]	1	Status of power shelf PSU 1
RF_PowerShelf_Status	total	Status of power shelf PSU, total
RF_Power_Supply_Energy [†]	1	Energy consumption of PSU 1 since boot
RF_Power_Supply_Energy	total	Energy consumption of all PSUs of power shelf since boot, total
RF_Power_Supply_FanSpeed [†]	1	Fan speed of PSU 1
RF_Power_Supply_FanSpeed	average	Fan speed of all PSUs of power shelf, averaged
RF_Power_Supply_Id [†]	1	
RF_Power_Supply_InputCurrent [†]	1	Input current for PSU 1
RF_Power_Supply_InputCurrent	total	Input current for all PSUs of power shelf, total

...continues

Table 9.4: Power shelf measurables...continued

Power Shelf Measurable	Parameter	Description
RF_Power_Supply_InputPower [†]	1	Input power for PSU 1
RF_Power_Supply_InputPower	total	Input power for all PSUs of power shelf, total
RF_Power_Supply_InputVoltage [†]	1	Input voltage for PSU 1
RF_Power_Supply_LifetimeReading_Energy [†]	1	Input lifetime reading for energy for PSU 1
RF_Power_Supply_Name [†]	1	
RF_Power_Supply_OutputPower [†]	1	Output power for PSU 1
RF_Power_Supply_OutputPower	total	Output power for all PSUs of power shelf, total
RF_Power_Supply_PowerFactor_Input [†]	1	
RF_Power_Supply_RailCurrent [†]	1	Rail Current for PSU 1
RF_Power_Supply_RailCurrent	total	Rail Current for all PSUs of power shelf, total
RF_Power_Supply_RailPower [†]	1	Rail power for PSU 1
RF_Power_Supply_RailPower	total	Rail power for all PSUs of power shelf, total
RF_Power_Supply_RailVoltage [†]	1	Rail voltage for PSUs 1
RF_Power_Supply_Temperature [†]	1	Temperature for PSU 1
RF_Power_Supply_Temperature	average	Temperature for all PSUs of power shelf, averaged
RF_Power_Supply_code_error [†]	1	Code error for PSU 1
RF_Power_Supply_message_error [†]	1	Message error for PSU 1
RF_Summary_Metrics_Id		
RF_Summary_Metrics_Name		
RF_Summary_Metrics_Power		Power consumption (Watts)
RF_Summary_Metrics_Temperature		Temperature (Celsius)
RF_Total_PSU		
TotalPowerShelfActivePSU*		Total power shelf PSUs active
TotalPowerShelfAvailablePSU*		Total power shelf PSUs available
TotalPowerShelfCriticalPSU‡		Are all active PSUs below critical threshold?
TotalPowerShelfDegradedPSU‡		Are all active PSUs below degraded threshold?
TotalPowerShelfHealthyPSU‡		Are all PSUs active?
TotalPowerShelfInputCurrent*		Total power shelf input current
TotalPowerShelfInputPower*		Total power shelf input power
TotalPowerShelfRailCurrent*		Total power shelf rail current
TotalPowerShelfRailPower*		Total power shelf rail power
TotalPowerShelfPSU*		Total power shelf PSU

* From AggregatePowerShelf data producer

** From ClusterTotal data producer

† Takes the PSU number as a parameter. On the DGX GB200 it can be from 1 to 6

‡ Health check

9.5 Cooling Distribution Unit Metrics On The DGX GB200

Metrics for the cooling distribution units (CDUs) of the DGX GB200 platform are shown in table 9.5. The data producers for these are:

- MonitoringSystem for the metrics in the table beginning with the string CDU
- AggregateCDU for the metrics in the table that begin with the strings Average or Total

Table 9.5: DGX GB200 metrics for CDUs

DGX GB200 CDU Metric	Description
AverageCDULiquidDifferentialPressure	Average CDU liquid cooling differential pressure (Pa)
AverageCDULiquidFlow	Average CDU liquid cooling flow (LPM)
AverageCDULiquidReturnTemperature	Average CDU liquid cooling return temperature (°C)
AverageCDULiquidSupplyTemperature	Average CDU liquid cooling supply temperature (°C)
AverageCDULiquidSystemPressure	Average CDU liquid system pressure (Pa)
CDUAvailable	CDU available
CDULiquidDifferentialPressure	CDU liquid differential pressure (Pa)
CDULiquidFlow	CDU liquid flow (LPM)
CDULiquidReturnTemperature	CDU liquid return temperature (°C)
CDULiquidSupplyTemperature	CDU liquid supply temperature (°C)
CDULiquidSystemPressure	CDU liquid system pressure (Pa)
CDUStatus	CDU status (l/s)
TotalCDUAvailable	Total CDU available
TotalCDUStatus	Total CDU status

9.6 GPU Measurables For The DGX GB200

Some basic GPU measurables are covered in appendix G of the *Administrator Manual*.

Extra GPU measurables available on the DGX GB200 platform are:

Table 9.6:

GPU Measurables For The DGX GB200	Parameter	Description
gpu_board_limitViolation*	gpu0	Board violation limit
gpu_board_limitViolation	total	Throttling duration due to board limits
gpu_c2c_link_bandwidth*	gpu0	GPU C2C link bandwidth
gpu_c2c_link_count*	gpu0	GPU C2C link count
gpu_c2c_link_status*	gpu0	GPU C2C link status
gpu_correctable_remapped_rows*	gpu0	Number of remapped rows for correctable errors
gpu_dec_utilization	average	Average GPU decoder utilization
gpu_dec_utilization	gpu0	GPU decoding usage
gpu_ecc_dbe_agg*	gpu0	Total double bit aggregate ECC errors
gpu_ecc_dbe_vol*	gpu0	Total double bit volatile ECC errors
gpu_ecc_sbe_agg*	gpu0	Total single bit aggregate ECC errors

...continues

Table 9.6: ...continued

GPU Measurables For The DGX GB200	Parameter	Description
gpu_ecc_sbe_vol*	gpu0	Total single bit volatile ECC errors
gpu_enc_utilization	average	Average GPU encoder utilization
gpu_enc_utilization*	gpu0	GPU encoding usage
gpu_enforced_power_limit*	gpu0	GPU enforced power limit
gpu_enforced_power_profile_mask*	gpu0	Enforced workload power profile mask
gpu_fabric_status*	gpu0	Status of the nvlink link domain (Enum)
gpu_fan_speed*	gpu0	GPU fan speed percentage
gpu_hbm_memory_temperature*	gpu0	High bandwidth memory (GPU memory) temperature
gpu_health_driver*,**	gpu0	Driver-related status
gpu_health_hostengine**		Host engine status
gpu_health_inforom*,**	gpu0	Inforom status
gpu_health_mcu*,**	gpu0	Microcontroller unit status
gpu_health_mem*,**	gpu0	Memory status
gpu_health_nvlink*,**	gpu0	NVLINK system status
gpu_health_nvswitch_fatal*,**	gpu0	NV switch fatal errors
gpu_health_nvswitch_non_fatal*,**	gpu0	NV switch non fatal errors
gpu_health_overall*		Overall system status
gpu_health_overall*,**	gpu0	Overall system status for the specified GPU
gpu_health_pcie*,**	gpu0	PCIe system status
gpu_health_pmu*,**	gpu0	Power management unit status
gpu_health_power*,**	gpu0	Power status
gpu_health_sm*,**	gpu0	Streaming multiprocessor status
gpu_health_thermal*,**	gpu0	Temperature status
gpu_low_utilViolation*	gpu0	Low utilization violation limit
gpu_low_utilViolation	total	Throttling duration due to low utilization limits
gpu_mem_clock*	gpu0	GPU memory clock
gpu_mem_copy_utilization	average	Average GPU memory copy utilization
gpu_mem_copy_utilization*	gpu0	Percentage of the GPU memory copy used
gpu_mem_free*	gpu0	Amount of GPU free memory
gpu_mem_total*	gpu0	GPU framebuffer size
gpu_mem_total	total	combined GPU total memory
gpu_mem_used*	gpu0	Amount of GPU used memory
gpu_mem_used	total	combined GPU memory usage
gpu_mem_utilization	average	Average GPU memory utilization
gpu_mem_utilization*	gpu0	Percentage of GPU memory used
gpu_memory_temp	average	Average GPU memory temperature
gpu_memory_temp*	gpu0	GPU memory temperature
gpu_nvlink_crc_data_errors*	gpu0	Total nvlink data CRC errors over all lanes
gpu_nvlink_crc_flit_errors*	gpu0	Total nvlink flit CRC errors over all lanes
gpu_nvlink_total_bandwidth*	gpu0	Total nvlink bandwidth used
gpu_nvlink_total_bandwidth	total	combined GPU nvlink bandwidth

...continues

Table 9.6: ...continued

GPU Measurables For The DGX GB200	Parameter	Description
gpu_performance_state*	gpu0	GPU performance state (0=highest)
gpu_power_management_limit*	gpu0	GPU power management limit
gpu_power_usage*	gpu0	GPU power usage
gpu_power_usage	total	combined GPU power usage
gpu_powerViolation*	gpu0	Throttling duration due to power constraints
gpu_powerViolation	total	Throttling duration due to power constraints
gpu_reliabilityViolation*	gpu0	Reliability violation limit
gpu_reliabilityViolation	total	Throttling duration due to reliability limits
gpu_requested_power_profile_mask*	gpu0	Requested workload power profile mask
gpu_row_remap_failure*	gpu0	Remapping of rows failures
gpu_shutdown_temp*	gpu0	GPU shutdown temperature
gpu_slowdown_temp*	gpu0	GPU slowdown temperature
gpu_sm_clock*	gpu0	GPU shader multiprocessor clock
gpu_sync_boostViolation*	gpu0	Throttling duration due to sync boost constraints
gpu_sync_boostViolation	total	Throttling duration due to sync boost constraints
gpu_temperature	average	Average GPU temperature
gpu_temperature*	gpu0	GPU temperature
gpu_thermalViolation*	gpu0	Throttling duration due to thermal constraints
gpu_thermalViolation	total	Throttling duration due to thermal constraints
gpu_total_app_clocksViolation*	gpu0	App clock violation limit
gpu_total_app_clocksViolation	total	Throttling duration due to application clocks limits
gpu_total_base_clocksViolation*	gpu0	Base clock violation limit
gpu_total_base_clocksViolation	total	Throttling duration due to base clocks limits
gpu_uncorrectable_remapped_rows*	gpu0	Number of remapped rows for uncorrectable errors
gpu_utilization	average	Average GPU utilization
gpu_utilization*	gpu0	GPU utilization percentage
gpu_xid_error*	gpu0	The value is the specific XID error

* parameter can be gpu0, gpu1...

** health check

9.7 Prometheus Metrics For The DGX GB200

Extra Prometheus metrics available on the DGX GB200 platform are shown in table 9.7.

The data producer for the metrics with a prefix of job_metadata_ is JobMetadataSampler. The remaining metrics are produced by JobSampler.

BCM picks up the measurables and their descriptions automatically. Many measurables still have no associated descriptions. For now, the names of the measurables can be a guide on what they are about.

Table 9.7:

Prometheus Metrics For The DGX GB200	Description
job_cpu_power_limit	CPU power limit
job_cpu_usage	
job_cpu_temperature	
job_cputacct_usage_seconds	Total CPU time consumed by processes
job_gpu_board_limitViolation	
job_gpu_c2c_link_count	
job_gpu_c2c_link_status	
job_gpu_correctable_remapped_rows	
job_gpu_dec_utilization	
job_gpu_ecc_dbe_agg	
job_gpu_ecc_dbe_vol	
job_gpu_ecc_sbe_agg	
job_gpu_ecc_sbe_vol	
job_gpu_enc_utilization	
job_gpu_enforced_power_limit	
job_gpu_enforced_power_profile_mask	
job_gpu_fabric_status	
job_gpu_fan_speed	
job_gpu_low_utilViolation	
job_gpu_mem_clock	
job_gpu_mem_copy_utilization	
job_gpu_mem_free	
job_gpu_mem_total	
job_gpu_mem_used	
job_gpu_mem_utilization	
job_gpu_memory_temp	
job_gpu_nvlink_crc_data_errors	
job_gpu_nvlink_crc_flit_errors	
job_gpu_nvlink_total_bandwidth	Total nvlink bandwidth used
job_gpu_performance_state	
job_gpu_power_management_limit	
job_gpu_power_usage	
job_gpu_powerViolation	
job_gpu_reliabilityViolation	
job_gpu_requested_power_profile_mask	
job_gpu_row_remap_failure	
job_gpu_shutdown_temp	
job_gpu_slowdown_temp	

...continues

Table 9.7: ...continued

Prometheus Metrics For The DGX GB200	Description
job_gpu_sm_clock	
job_gpu_sync_boostViolation	
job_gpu_temperature	
job_gpu_thermalViolation	
job_gpu_total_app_clocks_violation	
job_gpu_total_base_clocks_violation	
job_gpu_uncorrectable_remapped_rows	
job_gpu_utilization	
job_gpu_wasted	The amount of the allocated GPUs that is not used
job_gpu_workload_power_profile	GPU workload power profile status
job_gpu_xid_error	
job_memory_cache_bytes	Page cache, including tmpfs (shmem)
job_memory_failcnt	Number of times that the memory limit has reached the value set in memory.limit_in_bytes
job_memory_mapped_file_bytes	Size of memory-mapped mapped files, including tmpfs
job_memory_memsw_failcnt	Number of times that the memory plus swap space limit has reached the value set in memory.memsw.limit_in_bytes
job_memory_memsw_usage_bytes	Maximum amount of memory and swap space used by processes
job_memory_rss_bytes	Anonymous and swap cache, not including tmpfs
job_memory_swap_bytes	Swap memory usage
job_memory_usage_bytes	Total memory usage processes
job_metadata_allocated_cpu_cores	
job_metadata_allocated_gpus	
job_metadata_is_running	
job_metadata_is_waiting	
job_metadata_num_cpus	
job_metadata_num_nodes	
job_metadata_pending_jobs	
job_metadata_running_jobs	
job_metadata_running_seconds	
job_metadata_waiting_seconds	

9.8 Redfish Metrics For The DGX GB200

Redfish monitoring can be carried out for devices with a Redfish subscription. The `redfishsubscriptions` command can be used to check if a device has a subscription.

Example

```
[basecm11->device]# redfishsubscriptions node001
hostname      ...initialized ip          kind    port     ssl      subscribed...
-----  .....
node001       ...true           7.241.2.13  GB200   443     true     true
```

Redfish firmware metrics that may be available on the DGX GB200 platform are shown in table 9.8.

As is usual with hardware, the metrics can change according to hardware and especially with firmware changes. The table should therefore be regarded as an indication of what metrics are available, and not a list of what metrics must exist.

BCM picks up the metrics and their descriptions automatically. Many metrics have no associated descriptions. For now, the names of the metrics can be a guide on what they are about.

The list of Redfish metrics available can be found by running:

Example

```
basecm11->monitoring->measurable]# list metric | grep RF_
```

Table 9.8:

Redfish Metrics For The DGX GB200	Description
RF_AOC_NIC1Temp ^[1]	AOC NIC 1 Temperature
RF_AOC_PORTATemp	AOC PORT A Temperature
RF_AOC_PORTBTemp	AOC PORT B Temperature
RF_AOC_PORTCTemp	AOC PORT C Temperature
RF_Active_PSU ^[2]	
RF_Average_Temp	
RF_BF3_Slot_1_NIC_Temp_0	BF3 Slot 1 NIC Temperature 0
RF_BF3_Slot_2_NIC_Temp_0	BF3 Slot 2 NIC Temperature 0
RF_BF3_Slot_1_Port_0_Temp	BF3 Slot 1 Port 0 Temperature
RF_BF3_Slot_1_Port_1_Temp	BF3 Slot 1 Port 1 Temperature
RF_BF3_Slot_2_Port_0_Temp	BF3 Slot 2 Port 0 Temperature
RF_BF3_Slot_2_Port_1_Temp	BF3 Slot 2 Port 1 Temperature
RF_BF3_Slot_1_Temp	BF3 Slot 1 Temperature
RF_BF3_Slot_2_Temp	BF3 Slot 2 Temperature
RF_BMCTemp	BMC Temperature
RF_BMC_0_DCSCM_Temp_0	BMC 0 DCSCM Temperature 0
RF_C2CPU1Temp	C2 CPU1 Temperature
RF_C2CPU2Temp	C2 CPU2 Temperature
RF_CPU1RearTemp	CPU1 Rear Temperature
RF_CPU1Temp	CPU1 Temperature
RF_CPU1_VCCHV	CPU1_VCCHV

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_CPU1_VCCIN	CPU1_VCCIN
RF_CPU1_VCCON	CPU1_VCCON
RF_CPU1_VRMHVTemp	CPU1_VRMHV Temperature
RF_CPU1_VRMINTemp	CPU1_VRMIN Temperature
RF_CPU1_VRMONTemp	CPU1_VRMON Temperature
RF_CPU2RearTemp	CPU2 Rear Temperature
RF_CPU2Temp	CPU2 Temperature
RF_CPU2_VCCHV	CPU2_VCCHV
RF_CPU2_VCCIN	CPU2_VCCIN
RF_CPU2_VCCON	CPU2_VCCON
RF_CPU2_VRMHVTemp	CPU2_VRMHV Temperature
RF_CPU2_VRMINTemp	CPU2_VRMIN Temperature
RF_CPU2_VRMONTemp	CPU2_VRMON Temperature
RF_CPU_0_CoreUtil_0 ^[3]	CPU_0_CoreUtil_0
<i>In the preceding row, the substring CoreUtil_0 can also take values of CoreUtil_1, CoreUtil_2..., up to CoreUtil_71</i>	
RF_CPU_0_CpuFreq_0 ^[3]	CPU_0_CpuFreq_0
RF_CPU_0_Energy_0 ^[3]	CPU_0_Energy_0
RF_CPU_0_Power_0 ^[3]	CPU_0_Power_0
RF_CPU_0_Processor_code_error ^[3]	
RF_CPU_0_Processor_message_error ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_AccumulatedGPUContextUtilizationDuration ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_AccumulatedSMUtilizationDuration ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_EDPViolationState ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_GlobalSoftwareViolationThrottleDuration ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_HardwareViolationThrottleDuration ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_MemoryPageRetirement ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_MemorySpareChannelPresence ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_PCIERXBytes ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_PCIETXBytes ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_PerformanceState ^[3]	
RF_CPU_0_Processor_Metrics__Nvidia_	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
PowerBreakPerformanceState ^[3]	
RF_CPU_0_Processor_Metrics	
PowerLimitThrottleDuration ^[3]	
RF_CPU_0_Processor_Metrics_	
ThermalLimitThrottleDuration ^[3]	
RF_CPU_0_TempAvg_0 ^[3]	CPU_0_TempAvg_0
RF_CPU_0_TempLimit_0 ^[3]	CPU_0_TempLimit_0
RF_Chassis_0_FAN_1_FRONT ^[4]	Chassis 0 FAN 1 FRONT
RF_Chassis_0_FAN_1_PWM ^[4]	Chassis 0 FAN 1 PWM
RF_Chassis_0_FAN_1_REAR ^[4]	Chassis 0 FAN 1 REAR
RF_Chassis_0_Front_IO_Temp_	
0	Chassis 0 Front IO Temp 0
RF_Chassis_0_LeakDetector_	
0_ColdPlate ^[5]	Chassis 0 LeakDetector 0 ColdPlate
RF_Chassis_0_LeakDetector_	
0_Manifold ^[5]	Chassis 0 LeakDetector 0 Manifold
RF_Chassis_0_LeakDetector_	
1_ColdPlate ^[5]	Chassis 0 LeakDetector 1 ColdPlate
RF_Chassis_0_LeakDetector_	
1_Manifold ^[5]	Chassis 0 LeakDetector 1 Manifold
RF_Chassis_0_TotalHSC_Power_	
0	Chassis 0 TotalHSC Power 0
RF_Chassis_Intru_Reading	
RF_Chassis_Intru_ReadingRangeMax	
RF_Chassis_Intru_ReadingRangeMin	
RF_DIMM_SlotPartLocationContext_	
Location	
RF_DIMM_SlotState_Status	
RF_DIMM_Slot_BusWidthBits	
RF_DIMM_Slot_Capacity	
RF_DIMM_Slotcode_Error	
RF_DIMM_Slotmessage_Error	
RF_DIMM_Slot_OperatingSpeed	
RF_DIMM_Slot_DataWidthBits	
RF_DIMM_Slot_FirmwareRevision	
RF_DIMM_Slot_Nvidia_RowRemappingFailed	
RF_DIMM_Slot_Nvidia_RowRemappingPending	
RF_DIMM_Slot_OperatingSpeed	
RF_DIMM_Slot_Rank	
RF_Energy_0	Energy_0
RF_FPGATemp	FPGA Temp

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_FPGA_ARTTemp	FPGA_ART Temp
RF_GPU_0_DRAM_0_Memory_MetricsCorrectableE\ CCErrorCount_LifeTime ^[6]	
RF_GPU_0_DRAM_0_Memory_MetricsUncorrectabl\ eECCErrorCount_LifeTime ^[6]	
RF_GPU_0_DRAM_0_Memory_Metrics_ BandwidthUtilization ^[6]	
RF_GPU_0_DRAM_0_Memory_Metrics_ CapacityUtilizationPercent ^[6]	
RF_GPU_0_DRAM_0_Memory_Metrics_ OperatingSpeed ^[6]	
RF_GPU_0_Processor_MetricsCorrectableErrorCo\ unt_PCIEErrors ^[6]	
RF_GPU_0_Processor_MetricsFatalErrorCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsL0ToRecoveryCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsNAKReceivedCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsNAKSentCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsNonFatalErrorCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsReplayCount_ PCIeErrors ^[6]	
RF_GPU_0_Processor_MetricsReplayRolloverCoun\ t_PCIEErrors ^[6]	
RF_GPU_0_Processor_MetricsUnsupportedRequest\ Count_PCIEErrors ^[6]	
RF_GPU_0_Processor_Metrics_ BandwidthUtilization ^[6]	
RF_GPU_0_Processor_Metrics_ CoreVoltage ^[6]	
RF_GPU_0_Processor_Metrics_ LifeTime_CorrectableECCError ^[6]	
RF_GPU_0_Processor_Metrics_ LifeTime_UncorrectableECCError ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_ AccumulatedGPUContextUtilizationDuration ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_ AccumulatedSMUtilizationDuration ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_ DMMAUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_ FP16ActivityPercent ^[6]	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_GPU_0_Processor_Metrics_Nvidia_FP32ActivityPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_FP64ActivityPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_GlobalSoftwareViolationThrottleDuration ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_GraphicsEngineActivityPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_HMMUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_HardwareViolationThrottleDuration ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_IMMAUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_IntegerActivityUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVDecUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVJpgUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVLinkDataRxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVLinkDataTxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVLinkRawRxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVLinkRawTxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_NVOfaUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_PCIERXBytes ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_PCIERawRxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_PCIERawTxBandwidth ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_PCIETXBytes ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_SMActivityPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_SMOccupancyPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_SMUtilizationPercent ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
SRAMECCErrorThresholdExceeded ^[6]	
RF_GPU_0_Processor_Metrics_Nvidia_TensorCoreActivityPercent ^[6]	
RF_GPU_0_Processor_Metrics_OperatingSpeed ^[6]	
RF_GPU_0_Processor_Metrics_PowerLimitThrottleDuration ^[6]	
RF_GPU_0_Processor_Metrics_ThermalLimitThrottleDuration ^[6]	
RF_GPU_0_Processor_Metricscode_error ^[6]	
RF_GPU_0_Processor_Metricsmessage_error ^[6]	
RF_GPU_0_Processor_Reset_Metrics_ConventionalResetEntry ^[6]	
RF_GPU_0_Processor_Reset_Metrics_ConventionalResetExit ^[6]	
RF_GPU_0_Processor_Reset_Metrics_FundamentalResetEntry ^[6]	
RF_GPU_0_Processor_Reset_Metrics_FundamentalResetExit ^[6]	
RF_GPU_0_Processor_Reset_Metrics_IRoTResetExit ^[6]	
RF_GPU_0_Processor_Reset_Metrics_PF_FLR_ResetEntry ^[6]	
RF_GPU_0_Processor_Reset_Metrics_PF_FLR_ResetExit ^[6]	
RF_HGX_BMC_0_Temp_0	HGX BMC 0 Temp 0
RF_HGX_Baseboard_0CoreCount_ProcessorSummary	
RF_HGX_Baseboard_0Count_ProcessorSummary	
RF_HGX_Baseboard_0Model_ProcessorSummary	
RF_HGX_Baseboard_0State_Status	
RF_HGX_Baseboard_0TotalSystemMemoryGiB_MemorySummary	
RF_HGX_Baseboard_0_Description	
RF_HGX_Baseboard_0_LastResetTime	
RF_HGX_Baseboard_0_Nvidia_ISTMModeEnabled	
RF_HGX_Baseboard_0_PowerMode	
RF_HGX_Baseboard_0_PowerState	
RF_HGX_CPU_0State_Status ^[3]	
RF_HGX_CPU_0_AssetTag ^[3]	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_HGX_CPU_0_PartLocation_ServiceLabel ^[3]	
RF_HGX_CPU_0_SKU ^[3]	
RF_HGX_CPU_0code_error ^[3]	
RF_HGX_CPU_0message_error ^[3]	
RF_HGX_Chassis_0_TotalGPU_Power_0	HGX Chassis 0 TotalGPU Power 0
RF_HGX_GPU_0State_Status ^[6]	
RF_HGX_GPU_0_DRAM_0_Power_0 ^[6]	HGX GPU 0 DRAM 0 Power 0
RF_HGX_GPU_0_DRAM_0_Temp_0 ^[6]	HGX GPU 0 DRAM 0 Temp 0
RF_HGX_GPU_0_Energy_0 ^[6]	HGX GPU 0 Energy 0
RF_HGX_GPU_0_MaxPower ^[6]	
RF_HGX_GPU_0_MinPower ^[6]	
RF_HGX_GPU_0_Power_0 ^[6]	HGX GPU 0 Power 0
RF_HGX_GPU_0_SKU ^[6]	
RF_HGX_GPU_0_TEMP_0 ^[6]	HGX GPU 0 TEMP 0
RF_HGX_GPU_0_TEMP_1 ^[6]	HGX GPU 0 TEMP 1
RF_HGX_GPU_0_Voltage_0 ^[6]	HGX GPU 0 Voltage 0
RF_HGX_ProcessorModule_0_Exhaust_Temp_0 ^[12]	HGX ProcessorModule 0 Exhaust Temp 0
RF_HGX_ProcessorModule_0_Inlet_Temp_0 ^[12]	HGX ProcessorModule 0 Inlet Temp 0
RF_HGX_ProcessorModule_0_Inlet_Temp_1 ^[12]	HGX ProcessorModule 0 Inlet Temp 0
RF_HotswapTemp	Hotswap Temp
RF_IO_Board_0_CX7_0_Port_0_Temp	IO Board 0 CX7 0 Port 0 Temp
RF_IO_Board_0_CX7_0_Temp	IO Board 0 CX7 0 Temp
RF_IO_Board_0_CX7_0_Temp_0	IO Board 0 CX7 0 Temp 0
RF_IO_Board_0_CX7_1_Port_0_Temp	IO Board 0 CX7 1 Port 0 Temp
RF_IO_Board_0_CX7_1_Temp	IO Board 0 CX7 1 Temp
RF_IO_Board_0_CX7_1_Temp_0	IO Board 0 CX7 1 Temp 0
RF_IO_Board_1_CX7_0_Port_0_Temp	IO Board 1 CX7 0 Port 0 Temp
RF_IO_Board_1_CX7_0_Temp	IO Board 1 CX7 0 Temp
RF_IO_Board_1_CX7_0_Temp_0	IO Board 1 CX7 0 Temp 0

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_I0_Board_1_CX7_1_Port_0_Temp	IO Board 1 CX7 1 Port 0 Temp
RF_I0_Board_1_CX7_1_Temp	IO Board 1 CX7 1 Temp
RF_I0_Board_1_CX7_1_Temp_0	IO Board 1 CX7 1 Temp 0
RF_InletTemp	Inlet Temp
RF_InterswitchPort_0_Port_MetricsRXFrames_Networking	
RF_InterswitchPort_0_Port_MetricsRXMulticastFrames_Networking	
RF_InterswitchPort_0_Port_MetricsRXUnicastFrames_Networking	
RF_InterswitchPort_0_Port_MetricsTXDiscards_Networking	
RF_InterswitchPort_0_Port_MetricsTXFrames_Networking	
RF_InterswitchPort_0_Port_MetricsTXMulticastFrames_Networking	
RF_InterswitchPort_0_Port_MetricsTXUnicastFrames_Networking	
RF_InterswitchPort_0_Port_Metrics_Nvidia_BitErrorRate	
RF_InterswitchPort_0_Port_Metrics_Nvidia_LinkDowned	
RF_InterswitchPort_0_Port_Metrics_Nvidia_LinkErrorRecovery	
RF_InterswitchPort_0_Port_Metrics_Nvidia_MalformedPackets	
RF_InterswitchPort_0_Port_Metrics_Nvidia_NeighborMTUDiscards	
RF_InterswitchPort_0_Port_Metrics_Nvidia_QP1Dropped	
RF_InterswitchPort_0_Port_Metrics_Nvidia_RXRemotePhysicalErrors	
RF_InterswitchPort_0_Port_Metrics_Nvidia_RXSwitchRelayErrors	
RF_InterswitchPort_0_Port_Metrics_Nvidia_TXWait	
RF_InterswitchPort_0_Port_Metrics_Nvidia_VL15Dropped	
RF_InterswitchPort_0_Port_Metrics_Nvidia_VL15TBytes	
RF_InterswitchPort_0_Port_Metrics_Nvidia_VL15TPackets	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_InterSwitchPort_0_Port_Metrics_RXBytes	
RF_InterSwitchPort_0_Port_Metrics_RXErrors	
RF_InterSwitchPort_0_Port_Metrics_TXBytes	
RF_InterSwitchPort_0_ResourceState_Status	
RF_InterSwitchPort_0_Resource_CurrentSpeed	
RF_InterSwitchPort_0_Resource_LinkState	
RF_InterSwitchPort_0_Resource_LinkStatus	
RF_InterSwitchPort_0_Resource_MaxSpeed	
RF_LD_Chassis_0_LeakDetector_0_ColdPlate ^[5]	Chassis 0 LeakDetector 0 ColdPlate
RF_LD_Chassis_0_LeakDetector_0_Manifold ^[5]	Chassis 0 LeakDetector 0 Manifold
RF_LD_Chassis_0_LeakDetector_1_ColdPlate ^[5]	Chassis 0 LeakDetector 1 ColdPlate
RF_LD_Chassis_0_LeakDetector_1_Manifold ^[5]	Chassis 0 LeakDetector 1 Manifold
RF_LD_LeakDetection ^[5]	Leak Detection Systems
RF_LeakDetection ^[5]	Leak Detection Systems
RF_M2_SSD1Temp	M2_SSD1 Temp
RF_M2_SSD2Temp	M2_SSD2 Temp
RF_MB1.8VCC	MB 1.8VCC
RF_MB1.8VSB	MB 1.8VSB
RF_MB12V	MB 12V
RF_MB12VSB	MB 12VSB
RF_MB3.3VCC	MB 3.3VCC
RF_MB3.3VSB	MB 3.3VSB
RF_MB5VCC	MB 5VCC
RF_MB5VSB	MB 5VSB
RF_MGX_NVSwitch_0_TEMP_0	MGX NVSwitch 0 TEMP 0
RF_MGX_NVSwitch_1_TEMP_0	MGX NVSwitch 1 TEMP 0
RF_MemCtl_0_Freq_0	MemCtl_0_Freq_0
RF_MemCtl_1_Freq_0	MemCtl_1_Freq_0
RF_NVLinkManagement_0_Port_MetricsRXFrames_Networking	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_NVLinkManagement_0_Port_MetricsRXMulticastFrames_Networking	
RF_NVLinkManagement_0_Port_MetricsRXUnicastFrames_Networking	
RF_NVLinkManagement_0_Port_MetricsTXDiscards_Networking	
RF_NVLinkManagement_0_Port_MetricsTXFrames_Networking	
RF_NVLinkManagement_0_Port_MetricsTXMulticastFrames_Networking	
RF_NVLinkManagement_0_Port_MetricsTXUnicastFrames_Networking	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_BitErrorRate	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_LinkDowned	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_LinkErrorRecovery	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_MalformedPackets	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_NeighborMTUDiscards	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_QP1Dropped	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_RXRemotePhysicalErrors	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_RXSwitchRelayErrors	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_TXWait	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_VL15Dropped	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_VL15TXBytes	
RF_NVLinkManagement_0_Port_Metrics_Nvidia_VL15TXPackets	
RF_NVLinkManagement_0_Port_Metrics_RXBytes	
RF_NVLinkManagement_0_Port_Metrics_RXErrors	
RF_NVLinkManagement_0_Port_Metrics_TXBytes	
RF_NVLinkManagement_0_ResourceState_Status	
RF_NVLinkManagement_0_Resource_	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
CurrentSpeed	
RF_NVLinkManagement_0_Resource_LinkState	
RF_NVLinkManagement_0_Resource_LinkStatus	
RF_NVLinkManagement_0_Resource_MaxSpeed	
RF_NVLink_0_Port_MetricsRXFrames_Networking ^[7]	
RF_NVLink_0_Port_MetricsRXMulticastFrames_Networking ^[7]	
RF_NVLink_0_Port_MetricsRXUnicastFrames_Networking ^[7]	
RF_NVLink_0_Port_MetricsTXDiscards_Networking ^[7]	
RF_NVLink_0_Port_MetricsTxFrames_Networking ^[7]	
RF_NVLink_0_Port_MetricsTXMulticastFrames_Networking ^[7]	
RF_NVLink_0_Port_MetricsTXUnicastFrames_Networking ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_BitErrorRate ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_EffectiveError ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_LinkDowned ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_LinkErrorRecovery ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_MalformedPackets ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_NVLinkDataRxBandwidth ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_NVLinkDataTxBandwidth ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_NVLinkRawRxBandwidth ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_NVLinkRawTxBandwidth ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_NeighborMTUDiscards ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_QP1Dropped ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_RXNoProtocolBytes ^[7]	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_NVLink_0_Port_Metrics_Nvidia_RXRemotePhysicalErrors ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_RXSwitchRelayErrors ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_TXNoProtocolBytes ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_TXWait ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_VL15Dropped ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_VL15TXBytes ^[7]	
RF_NVLink_0_Port_Metrics_Nvidia_VL15TXPackets ^[7]	
RF_NVLink_0_Port_Metrics_RXBytes ^[7]	
RF_NVLink_0_Port_Metrics_RXErrors ^[7]	
RF_NVLink_0_Port_Metrics_TXBytes ^[7]	
RF_NVLink_0_ResourceState_Status ^[7]	
RF_NVLink_0_Resource_CurrentSpeed ^[7]	
RF_NVLink_0_Resource_LinkState ^[7]	
RF_NVLink_0_Resource_LinkStatus ^[7]	
RF_NVLink_0_Resource_MaxSpeed ^[7]	
RF_NVLink_0_Resource_Nvidia_RXWidth ^[7]	
RF_NVLink_0_Resource_Nvidia_TXWidth ^[7]	
RF_NVME_M2_0_Temp_0	NVME M2 0 Temp 0
RF_NVMe_SSDATemp	NVMe_SSDA Temp
RF_NVSwitch_0_ResourceState_Status	
RF_NVSwitch_0_Resource_#Switch.Reset_target	
RF_NVSwitch_0_Resource_CurrentBandwidth	
RF_NVSwitch_0_Resource_Enabled	
RF_NVSwitch_0_Resource_FirmwareVersion	
RF_NVSwitch_0_Resource_MaxBandwidth	
RF_NVSwitch_0_Resource_Nvidia_SwitchIsolationMode	
RF_NVSwitch_1_ResourceState_	

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
Status	
RF_NVSwitch_1_Resource_#Switch.Reset_target	
RF_NVSwitch_1_Resource_CurrentBandwidth	
RF_NVSwitch_1_Resource_Enabled	
RF_NVSwitch_1_Resource_FirmwareVersion	
RF_NVSwitch_1_Resource_MaxBandwidth	
RF_NVSwitch_1_Resource_Nvidia_SwitchIsolationMode	
RF_P1_DIMMA_DTemp	P1_DIMMA D Temp
RF_P1_DIMME_HTemp	P1_DIMME H Temp
RF_P2_DIMMA_DTemp	P2_DIMMA D Temp
RF_P2_DIMME_HTemp	P2_DIMME H Temp
RF_PCB_Inlet_Temp	PCB Inlet Temp
RF_PCH1.05V	PCH 1.05V
RF_PCHPVNN	PCH PVNN
RF_PCHTemp	PCH Temp
RF_PCIE_0_ResourceState_Status	
RF_PCIE_0_Resource_ActiveWidth	
RF_PCIE_0_Resource_CurrentSpeed	
RF_PCIE_0_Resource_MaxSpeed	
RF_PCIE_0_Resource_Width	
RF_PCIE_DeviceLanesInUse_PCIEInterface	
RF_PCIE_DeviceMaxLanes_PCIEInterface	
RF_PCIE_DeviceMaxPCIeType_PCIEInterface	
RF_PCIE_DevicePCIeType_PCIEInterface	
RF_PCIE_DeviceState_Status	
RF_PCIE_Device_Nvidia_AERCorrectableErrorStatus	
RF_PCIE_Device_Nvidia_AERUncorrectableErrorStatus	
RF_PCIE_Device_Nvidia_NVLinkReferenceClockEnabled	
RF_PCIE_Device_Nvidia_PCIEReferenceClockEnabled	
RF_PCIE_Devicecode_error	
RF_PCIE_Devicemessage_error	
RF_PDB_0_HSC_0_Cur_0	PDB 0 HSC 0 Cur 0
RF_PDB_0_HSC_0_Pwr_0	PDB 0 HSC 0 Pwr 0
RF_PDB_0_HSC_0_Temp_0	PDB 0 HSC 0 Temp 0
RF_PDB_0_HSC_0_Volt_In_0	PDB 0 HSC 0 Volt In 0
RF_PDB_0_HSC_0_Volt_Out_0	PDB 0 HSC 0 Volt Out 0

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_PDB_0_HSC_1_Cur_0	PDB 0 HSC 1 Cur 0
RF_PDB_0_HSC_1_Pwr_0	PDB 0 HSC 1 Pwr 0
RF_PDB_0_HSC_1_Temp_0	PDB 0 HSC 1 Temp 0
RF_PDB_0_HSC_1_Volt_In_0	PDB 0 HSC 1 Volt In 0
RF_PDB_0_HSC_1_Volt_Out_0	PDB 0 HSC 1 Volt Out 0
RF_PDB_0_Inlet_Temp_0	PDB 0 Inlet Temp 0
RF_PDB_0_Vreg_0_Temp_0	PDB 0 Vreg 0 Temp 0
RF_PDB_0_Vreg_0_Temp_1	PDB 0 Vreg 0 Temp 1
RF_PDB_0_Vreg_1_Temp_0	PDB 0 Vreg 1 Temp 0
RF_PDB_0_Vreg_1_Temp_1	PDB 0 Vreg 1 Temp 1
RF_PDB_TOTAL_CURRENT	
RF_PDB_TOTAL_PWR	
RF_PDB_TOTAL_VOLT_IN	
RF_PDB_TOTAL_VOLT_OUT	
RF_PS1_Status_Reading	
RF_PS1_Status_ReadingRangeMax	
RF_PS1_Status_ReadingRangeMin	
RF_PWConsumption	PW Consumption
RF_PeripheralTemp	Peripheral Temp
RF_PowerShelf_Status ^{[2],[8]}	Power shelf status of PSU
RF_PowerShelf_Status ^{[2],[8]}	Power shelf status total
RF_Power_0	Power_0
RF_Power_PowerControl	
RF_Power_PowerSupplies	
RF_Power_Redundancy	
RF_Power_Supply_Energy ^{[2],[8]}	Energy consumption since boot of PSU
RF_Power_Supply_Energy ^[9]	Energy consumption since boot of all PSUs
RF_Power_Supply_FanSpeed ^{[2],[8]}	Fan speed of PSU
RF_Power_Supply_FanSpeed ^[10]	Fan speed average of all PSUs
RF_Power_Supply_Id ^{[2],[8]}	ID of PSU
RF_Power_Supply_InputCurrent ^{[2],[8]}	Input current for PSU
RF_Power_Supply_InputCurrent ^[9]	Input current of all PSUs
RF_Power_Supply_InputPower ^{[2],[8]}	Input power for PSU
RF_Power_Supply_InputPower ^[9]	Input power total of all PSUs
RF_Power_Supply_InputVoltage ^{[2],[8]}	Input voltage for PSU
RF_Power_Supply_LifetimeReading_Energy ^{[2],[8]}	Energy consumption of PSU over lifetime

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_Power_Supply_Name ^{[2],[8]}	Name of PSU
RF_Power_Supply_OutputPower ^{[2],[8]}	Output power of PSU
RF_Power_Supply_OutputPower ^[9]	Output power total of all PSUs
RF_Power_Supply_PowerFactor_Input ^{[2],[8]}	Input power factor of PSU
RF_Power_Supply_RailCurrent ^{[2],[8]}	Rail current of PSU
RF_Power_Supply_RailCurrent ^[9]	Rail current total of all PSUs
RF_Power_Supply_RailPower ^{[2],[8]}	Rail current of PSU
RF_Power_Supply_RailPower ^[9]	Rail power total of all PSUs
RF_Power_Supply_RailVoltage ^{[2],[8]}	Rail voltage of PSU
RF_Power_Supply_Temperature ^{[2],[8]}	Temperature of PSU
RF_Power_Supply_Temperature ^[10]	Temperature average of all PSUs
RF_Power_Supply_code_error ^{[2],[8]}	Code error of PSU
RF_Power_Supply_message_error ^{[2],[8]}	Message error of PSU
RF_Power_Voltages	
RF_ProcessorModule_0_CPU_0_CoreUtil_0 ^[11]	ProcessorModule 0 CPU 0 CoreUtil 0. The substring CoreUtil_0 can also take values of CoreUtil_1, CoreUtil_2... up to CoreUtil_71
RF_ProcessorModule_0_CPU_0_CpuFreq_0 ^[11]	ProcessorModule 0 CPU 0 CpuFreq 0
RF_ProcessorModule_0_CPU_0_Energy_0 ^[11]	ProcessorModule 0 CPU 0 Energy 0
RF_ProcessorModule_0_CPU_0_EnforcedEDPc_0 ^[11]	ProcessorModule 0 CPU 0 EnforcedEDPc 0
RF_ProcessorModule_0_CPU_0_EnforcedEDPp_0 ^[11]	ProcessorModule 0 CPU 0 EnforcedEDPp 0
RF_ProcessorModule_0_CPU_0_Power_0 ^[11]	ProcessorModule 0 CPU 0 Power 0
RF_ProcessorModule_0_CPU_0_TempAvg_0 ^[11]	ProcessorModule 0 CPU 0 TempAvg 0
RF_ProcessorModule_0_CPU_0_TempLimit_0 ^[11]	ProcessorModule 0 CPU 0 TempLimit 0
RF_ProcessorModule_0_MemCntl_0_Freq_0 ^[11]	ProcessorModule 0 MemCntl 0 Freq 0
RF_ProcessorModule_0_Vreg_0_CpuPower_0 ^[11]	ProcessorModule 0 Vreg 0 CpuPower 0
RF_ProcessorModule_0_Vreg_0_CpuVoltage_0 ^[11]	ProcessorModule 0 Vreg 0 CpuVoltage 0
RF_ProcessorModule_0_Vreg_0_SocPower_0 ^[11]	ProcessorModule 0 Vreg 0 SocPower 0
RF_ProcessorModule_0_Vreg_0_SocVoltage_0 ^[11]	ProcessorModule 0 Vreg 0 SocVoltage 0

...continues

Table 9.8: ...continued

Redfish Metrics For The DGX GB200	Description
RF_SYS_ART1Temp	SYS_ART1 Temp
RF_SYS_ART2Temp	SYS_ART2 Temp
RF_StorageBackplane_0_SSD_0_Temp_0	StorageBackplane 0 SSD 0 Temp 0
RF_StorageBackplane_0_SSD_2_Temp_0	StorageBackplane 0 SSD 2 Temp 0
RF_StorageBackplane_1_SSD_0_Temp_0	StorageBackplane 1 SSD 0 Temp 0
RF_StorageBackplane_1_SSD_2_Temp_0	StorageBackplane 1 SSD 2 Temp 0
RF_Summary_Metrics_Id ^[2]	
RF_Summary_Metrics_Name ^[2]	
RF_Summary_Metrics_Power ^[2]	Power consumption (Watts)
RF_Summary_Metrics_Temperature ^[2]	Temperature (Celsius)
RF_SystemTemp	System Temp
RF_Total_PSU ^[2]	
RF_VBAT_Reading	
RF_VBAT_ReadingRangeMax	
RF_VBAT_ReadingRangeMin	
RF_Vreg_0_CpuPower_0	Vreg_0_CpuPower_0
RF_Vreg_0_CpuVoltage_0	Vreg_0_CpuVoltage_0
RF_Vreg_0_SocPower_0	Vreg_0_SocPower_0
RF_Vreg_0_SocVoltage_0	Vreg_0_SocVoltage_0
RF_Vreg_1_CpuPower_0	Vreg_1_CpuPower_0
RF_Vreg_1_CpuVoltage_0	Vreg_1_CpuVoltage_0
RF_Vreg_1_SocPower_0	Vreg_1_SocPower_0
RF_Vreg_1_SocVoltage_0	Vreg_1_SocVoltage_0

[1] The substring NIC1Temp can also be NIC2Temp, NIC3Temp...

[2] from redfish_power_shelf data producer

[3] The substring CPU_0 can also be CPU_1

[4] The substring FAN_1 can also be FAN_2, FAN_3...

[5] From redfish_leak data producer

[6] The substring GPU_0 can also be GPU_1, GPU_2...

[7] On the DGX GB200 the value of <number> in ...NVLink_<number>... can be from 0 to 71

[8] Takes the PSU number as a parameter. On the DGX GB200 it can be from 1 to 6

[9] Takes total as a parameter. On the DGX GB200 it means the total value of the 6 PSUs

[10] Takes average as a parameter. On the DGX GB200 it means the average value of the 6 PSUs

[11] On the DGX GB200 the value of <number> in ...ProcessorModule_<number>... can be 0 or 1

[12] The substring ProcessorModule_0 can also be ProcessorModule_1

9.9 Redfish Enums For The DGX GB200

Redfish monitoring can be carried out for devices with a Redfish subscription (page 70).

Redfish firmware enums that may be available on the DGX GB200 platform are shown in table 9.9.

As is usual with hardware, the metrics can change according to hardware and especially with firmware

changes. The table should therefore be regarded as an indication of what metrics are available, and not a list of what metrics must exist.

The list of Redfish enums that are there on the DGX GB200 platform can be found by running:

Example

```
basecm11->monitoring->measurable]% list enum | grep RF_
```

Table 9.9:

Redfish Enums For The DGX GB200	Parameter	Description
RF_BF3_Slot_1_NIC_Temp_0	state	
RF_BF3_Slot_2_NIC_Temp_0	state	
RF_BF3_Slot_2_Port_0_Temp	state	
RF_BF3_Slot_2_Port_1_Temp	state	
RF_BF3_Slot_2_Temp	state	
RF_BMC_0_DCSCM_Temp_0	state	
RF_Chassis_0_FAN_1_FRONT ^[1]	state	
RF_Chassis_0_FAN_1_PWM	state	
RF_Chassis_0_FAN_1_REAR	state	
RF_CPU_0_Processor__Nvidia_EDPViolationState		
RF_CPU_0_Processor__Nvidia_PerformanceState		
RF_CPU_0_Processor__Nvidia_PowerBreakPerform\anceState		
RF_CPU_1_Processor__Nvidia_EDPViolationState		
RF_CPU_1_Processor__Nvidia_PerformanceState		
RF_CPU_1_Processor__Nvidia_PowerBreakPerform\anceState		
RF_DIMM_SlotHealthRollup_Status	state	
RF_DIMM_SlotHealth_Status	state	
RF_ERoT_BMC_0_PowerState		
RF_ERoT_BMC_0HealthRollup_Status	state	
RF_ERoT_BMC_0Health_Status	state	
RF_ERoT_BMC_0_PowerState		
RF_HGX_BMC_0_Temp_0	state	
RF_HGX_Baseboard_0HealthRollup_Status	state	
RF_HGX_Baseboard_0Health_Status	state	
RF_HGX_Baseboard_0_PowerState		
RF_HGX_CPU_0HealthRollup_Status	state	
RF_HGX_CPU_0Health_Status	state	
RF_HGX_CPU_0_PowerState		

...continues

Table 9.9: ...continued

Redfish Enums For The DGX GB200	Parameter	Description
RF_HGX_CPU_1HealthRollup_Status	state	
RF_HGX_CPU_1Health_Status	state	
RF_HGX_CPU_1_PowerState		
RF_HGX_Chassis_0_TotalGPU_Power_0	state	
RF_HGX_ERoT_BMC_0HealthRollup_Status	state	
RF_HGX_ERoT_BMC_0Health_Status	state	
RF_HGX_ERoT_CPU_0HealthRollup_Status	state	
RF_HGX_ERoT_CPU_0Health_Status	state	
RF_HGX_ERoT_CPU_0_PowerState		
RF_HGX_ERoT_CPU_1HealthRollup_Status	state	
RF_HGX_ERoT_CPU_1Health_Status	state	
RF_HGX_ERoT_FPGA_0HealthRollup_Status	state	
RF_HGX_ERoT_FPGA_0Health_Status	state	
RF_HGX_ERoT_FPGA_1HealthRollup_Status	state	
RF_HGX_ERoT_FPGA_1Health_Status	state	
RF_HGX_CPU_0_PowerState		
RF_HGX_CPU_1_PowerState		
RF_HGX_GPU_0HealthRollup_Status ^[2]	state	
RF_HGX_GPU_0Health_Status ^[2]	state	
RF_HGX_GPU_0_DRAM_0_Power_0 ^[2]	state	
RF_HGX_GPU_0_DRAM_0_Temp_0 ^[2]	state	
RF_HGX_GPU_0_Energy_0 ^[2]	state	
RF_HGX_GPU_0_PowerState ^[2]		
RF_HGX_GPU_0_Power_0 ^[2]	state	
RF_HGX_GPU_0_TEMP_0 ^[2]	state	
RF_HGX_GPU_0_TEMP_1 ^[2]	state	
RF_HGX_GPU_0_Voltage_0 ^[2]	state	
RF_HGX_IRoT_GPU_0HealthRollup_Status ^[2]	state	
RF_HGX_IRoT_GPU_0Health_Status ^[2]	state	
RF_HGX_IRoT_GPU_1_PowerState ^[2]		
RF_HGX_ProcessorModule_0_Exhaust_Temp_0	state	
RF_HGX_ProcessorModule_0_Inlet_Temp_0	state	
RF_HGX_ProcessorModule_0_Inlet_Temp_1	state	
RF_HGX_ProcessorModule_1_Exhaust_Temp_0	state	
RF_HGX_ProcessorModule_1_Inlet_Temp_0	state	
RF_HGX_ProcessorModule_1_Inlet_Temp_1	state	
RF_IO_Board_0_CX7_0_Port_0_Temp	state	
RF_IO_Board_0_CX7_0_Temp	state	
RF_IO_Board_0_CX7_0_Temp_0	state	

...continues

Table 9.9: ...continued

Redfish Enums For The DGX GB200	Parameter	Description
RF_I0_Board_0_CX7_1_Port_0_Temp	state	
RF_I0_Board_0_CX7_1_Temp	state	
RF_I0_Board_0_CX7_1_Temp_0	state	
RF_I0_Board_1_CX7_0_Port_0_Temp	state	
RF_I0_Board_1_CX7_0_Temp	state	
RF_I0_Board_1_CX7_0_Temp_0	state	
RF_I0_Board_1_CX7_1_Port_0_Temp	state	
RF_I0_Board_1_CX7_1_Temp	state	
RF_I0_Board_1_CX7_1_Temp_0	state	
RF_IRoT_CX7_0HealthRollup_Status	state	
RF_IRoT_CX7_0Health_Status	state	
RF_IRoT_CX7_1HealthRollup_Status	state	
RF_IRoT_CX7_1Health_Status	state	
RF_IRoT_CX7_2HealthRollup_Status	state	
RF_IRoT_CX7_2Health_Status	state	
RF_IRoT_CX7_3HealthRollup_Status	state	
RF_IRoT_CX7_3Health_Status	state	
RF_LD_LeakDetection ^[3]		
RF_LeakDetector ^[3]	Chassis_0_LeakDetector_0_ColdPlate	
RF_LeakDetector ^[3]	Chassis_0_LeakDetector_0_Manifold	
RF_LeakDetector ^[3]	Chassis_0_LeakDetector_1_ColdPlate	
RF_LeakDetector ^[3]	Chassis_0_LeakDetector_1_Manifold	
RF_NVLinkManagement_0_ResourceHealthRollup_Status	state	
RF_NVLinkManagement_0_ResourceHealth_Status	state	
RF_NVLink_ResourceHealthRollup_Status	state	
RF_NVLink_ResourceHealth_Status	state	
RF_NVLink_ResourceState_Status ^[4]	acp0	
RF_NVLink_Resource_LinkState ^[4]	acp0	
RF_NVLink_0_ResourceHealthRollup_Status	state	
RF_NVLink_0_ResourceHealth_Status	state	
RF_NVME_M2_0_Temp_0	state	
RF_NVSwitch_0_ResourceHealthRollup_Status ^[5]	state	

...continues

Table 9.9: ...continued

Redfish Enums For The DGX GB200	Parameter	Description
RF_NVSwitch_0_ResourceHealth_Status ^[5]	state	
RF_InterswitchPort_0_ResourceHealthRollup_Status	state	
RF_InterswitchPort_0_ResourceHealth_Status	state	
RF_PCIE_0_ResourceHealthRollup_Status	state	
RF_PCIE_0_ResourceHealth_Status	state	
RF_PCIE_DeviceHealthRollup_Status	state	
RF_PCIE_DeviceHealth_Status	state	
RF_PDB_0_HSC_0_Cur_0	state	
RF_PDB_0_HSC_0_Pwr_0	state	
RF_PDB_0_HSC_0_Temp_0	state	
RF_PDB_0_HSC_0_Volt_In_0	state	
RF_PDB_0_HSC_0_Volt_Out_0	state	
RF_PDB_0_HSC_1_Cur_0	state	
RF_PDB_0_HSC_1_Pwr_0	state	
RF_PDB_0_HSC_1_Temp_0	state	
RF_PDB_0_HSC_1_Volt_In_0	state	
RF_PDB_0_HSC_1_Volt_Out_0	state	
RF_PDB_0_Inlet_Temp_0	state	
RF_PDB_0_Vreg_0_Temp_0	state	
RF_PDB_0_Vreg_0_Temp_1	state	
RF_PDB_0_Vreg_1_Temp_0	state	
RF_PDB_0_Vreg_1_Temp_1	state	
RF_ProcessorModule_0_CPU_0_CpuFreq_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_Energy_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_EnforcedEDPc_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_EnforcedEDPp_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_Power_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_TempAvg_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_TempLimit_0 ^[6]	state	
RF_ProcessorModule_0_MemCntl_0_Freq_0 ^[6]	state	
RF_ProcessorModule_0_Vreg_0_CpuPower_0 ^[6]	state	
RF_ProcessorModule_0_Vreg_0_CpuVoltage_0 ^[6]	state	
RF_ProcessorModule_0_Vreg_0_SocPower_0 ^[6]	state	
RF_ProcessorModule_0_Vreg_0_SocVoltage_0 ^[6]	state	
RF_ProcessorModule_0_CPU_0_CoreUtil_0 ^[6]	state	
<i>In the preceding row, the substring CoreUtil_0 can also take values of CoreUtil_1, CoreUtil_2..., up to CoreUtil_71</i>		
RF_StorageBackplane_0_SSD_0_Temp_0 ^[7]	state	

...continues

Table 9.9: ...continued

Redfish Enums For The DGX GB200	Parameter	Description
[1]	The number used in the substring Chassis_0 and the substring FAN_ can vary	
[2]	The substring GPU_0 can be GPU_1, GPU_2...	
[3]	Produced by redfish_leak data producer	
[4]	The parameter substring acp0 can be acp1, acp2...	
[5]	The substring NVSwitch_0 can be NVSwitch_1, NVSwitch_2...	
[6]	The substring ProcessorModule_0 can be ProcessorModule_1, ProcessorModule_2...	
[7]	The substring SSD_0 can be SSD_1, SSD_2..., and the substring StorageBackplane_0 can be StorageBackplane_1, StorageBackplane_2...	

9.10 Health Checks For The DGX GB200

Redfish health checks that may be additionally available on the DGX GB200 platform are shown in table 9.10.

As is usual with hardware, the health checks can change according to hardware and especially with firmware changes. The table should therefore be regarded as an indication of the health checks that are available, and not a list of the health checks that must exist.

The list of Redfish health checks that are there can be found by running:

Example

```
[basecm11->monitoring->measurable]% list healthcheck
```

Table 9.10: DGX GB200 Health Checks

DGX GB200 Health Check	Parameter	Description
gpu_health_driver	gpu0	Driver-related status
gpu_health_hostengine		Host engine status
gpu_health_inforom	gpu0	Inforom status
gpu_health_mcu	gpu0	Microcontroller unit status
gpu_health_mem	gpu0	Memory status
gpu_health_nvlink	gpu0	NVLINK system status
gpu_health_nvswitch_fatal	gpu0	NV switch fatal errors
gpu_health_nvswitch_non_fatal	gpu0	NV switch non-fatal errors
gpu_health_overall		Overall system status
gpu_health_overall	gpu0	Overall system status for GPU0
gpu_health_pcie	gpu0	PCIe system status
gpu_health_pmu	gpu0	Power management unit status
gpu_health_power	gpu0	Power status
gpu_health_sm	gpu0	Streaming multiprocessor status
gpu_health_thermal	gpu0	Temperature status
nvsm_pci_health		Checks health of PCI devices in a DGX system

...continues

Table 9.10: DGX GB200 Health Checks...continued

DGX GB200 Health Check	Parameter	Description
nvsm_show_health		Checks health of DGX system
NMXController		Is NMX Controller in the correct state?